

NỀN TẢNG AI TẠO SINH
(IT5410 – Foundation of Generative AI)

MỘT SỐ VẤN ĐỀ CỦA LÝ THUYẾT HỌC MÁY VÀ HỌC SÂU

Thân Quang Khoát

Trường CNTT&TT, ĐHBKH

2025

Nội dung

- Mở đầu
- **Một số vấn đề của Học sâu**
- Một số kiến trúc mạng nơron
- Mô hình sinh sâu
- Đánh giá chất lượng
- Học tăng cường

Learning theory

Basic concepts

The supervised learning problem


- There is an **unknown** (measurable) function
(Có một hàm gán nhãn nào đó)

$$y^*: x \mapsto y$$

- It maps each input $x \in \mathcal{X}$ to a label (output) $y \in \mathcal{Y}$
- Sometimes known as the *labeling function*
- Spaces: *input space* \mathcal{X} , *output space* \mathcal{Y}

- We can collect a dataset $\mathbf{D} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_m, y_m)\}$

- $y_i = y^*(x_i)$ for any $i \in \{1, \dots, m\}$
- x_i is an observation about input x (một quan sát về x)
- We need to learn y^* from \mathbf{D}



Chúng ta chưa quan tâm đến trường hợp y^* bị **nhiễu** hoặc **ngẫu nhiên** (ở đầu vào hoặc đầu ra).



The sampling distribution

- Dataset $\mathbf{D} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_m, y_m)\}$
 - Each \mathbf{x}_i is collected/generated according to a distribution P , i.e., \mathbf{x}_i is a sample from distribution P
 - P is often **unknown** in practice.

- Those samples are often assumed to be i.i.d.

(Independent and identically distributed – Độc lập với nhau và được sinh ra từ một phân bố)

- Easy for analysis, each example may contain maximal information



The
data-generation
model

The sampling distribution

- Dataset $\mathbf{D} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_m, y_m)\}$
- **Hardness** of the learning problem: **The two unknown**
 - y^* in the space of all measurable functions
 - $P(\mathbf{x})$ in the space of all probability distributions
- In practice, we often find a function h to **approximate** y^*
 - h is often called **predictor** or *prediction rule*



Learning algorithm

- **Learning algorithm (Learner):** an algorithm that can return one predictor h , based on a given training set D

$$\mathcal{A}: D \mapsto h$$

- Receive a training set D ,
- Can access the input space \mathcal{X} and label space \mathcal{Y}
- Return one predictor $h: \mathcal{X} \rightarrow \mathcal{Y}$
- What space does a learning algorithm map from?
- E.g., K-nearest neighbors, ordinary least square
- The learner **may not know** any thing about the true y^* and $P(x)$

(Thuật toán học thường không biết gì thêm về hàm y^* và phân bố P)

Hypothesis space


- Usually, we can choose a restricted set \mathcal{H} of functions

- A learner must find one $h \in \mathcal{H}$
- Our hope: $h(x) \approx y^*(x), \quad \forall x \in \mathcal{X}$

- Hypothesis space (model space):

a set \mathcal{H} of functions, providing candidates h for a learning algorithm

- Each candidate depends on input features: $h: \mathcal{X} \rightarrow \mathcal{Y}$
- Represents prior knowledge about a task
- Represents our **inductive bias** or preference
- Each h is often called a “**model**” or “**predictor**”
- E.g., Linear model: $\mathcal{H} = \{h(x, w) = w_0 + w_1 x_1 + \dots + w_n x_n: w_0, w_1, \dots, w_n \in \mathbb{R}\}$



Loss function

- How to see the goodness of each $h \in \mathcal{H}$?


- How well does h predict each input \mathbf{x} ?
- How to compare two predictors h and g ?

- **Loss (cost) function:** $f: \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}$

- $f(y, \hat{y})$: the cost/loss of prediction \hat{y} about y
- Tells the quality of each prediction

- Examples:

- 0-1 loss: $f(y, \hat{y}) = \mathbf{1}_{y \neq \hat{y}}$
- Square loss: $f(y, \hat{y}) = (y - \hat{y})^2$
- Absolute loss: $f(y, \hat{y}) = |y - \hat{y}|$
- Hinge loss: $f(y, \hat{y}) = \max(0, 1 - y\hat{y})$



For vectors:

- Square (ℓ_2) loss: $f(\mathbf{y}, \hat{\mathbf{y}}) = \|\mathbf{y} - \hat{\mathbf{y}}\|_2^2$
- Absolute (ℓ_1) loss: $f(\mathbf{y}, \hat{\mathbf{y}}) = \|\mathbf{y} - \hat{\mathbf{y}}\|_1$

Loss or Risk

- When h makes a prediction about \mathbf{x} , its loss is $f(y^*(\mathbf{x}), h(\mathbf{x}))$
- Empirical loss (lỗi thực nghiệm): the loss of a predictor h on the training set \mathbf{D}

$$F(\mathbf{D}, h) = \frac{1}{M} \sum_{i=1}^M f(y_i, h(\mathbf{x}_i))$$


- Sometimes known as **Empirical risk**
- Expected loss: the loss over the whole space

$$F(P, h) = \mathbb{E}_{\mathbf{x} \sim P} [f(y^*(\mathbf{x}), h(\mathbf{x}))]$$

- Sometimes known as **Expected risk**
- For binary classification, where $\mathcal{Y} = \{0,1\}$, and absolute loss f :

$$\Pr(y^*(\mathbf{x}) \neq h(\mathbf{x})) = F(P, h)$$

- For $\mathcal{Y} = \{-1, 1\}$ and absolute loss f : $\Pr(y^*(\mathbf{x}) \neq h(\mathbf{x})) = \frac{1}{2}F(P, h)$





y^*

Learning goal

- Learner's task: find a good $h \in \mathcal{H}$, based on a training set \mathbf{D}
- *Learning goal:* find one $h \in \mathcal{H}$ with *small expected loss*
 - It should **generalize** well on future data instances
 - Small training loss/error is not enough
- Ultimately, we want to find the best one in \mathcal{H}

$$h^* = \arg \min_{h \in \mathcal{H}} F(P, h)$$

- **Learning \neq Fitting**
 - Fitting focuses on minimizing the training loss $F(\mathbf{D}, h)$




Empirical Risk Minimization (ERM)

- In practice, a learner cannot access $F(P, h)$
- It should rely on $F(\mathbf{D}, h)$
 - Can be computed from a training set \mathbf{D}
- ERM will find

$$h_{erm} = \arg \min_{h \in \mathcal{H}} F(\mathbf{D}, h)$$

- Focuses on fitting (minimizing the training error)
- However, h_{erm} may get overfitting
- ERM is good in the cases of large training sets, due to

$$F(\mathbf{D}, h) \xrightarrow{m \rightarrow \infty} F(P, h)$$



- Regularization-based algorithms:

$$h_{reg} = \arg \min_{h \in \mathcal{H}} F(\mathbf{D}, h) + \lambda \cdot g(h)$$


- $g(h)$ is a penalty on h , λ is a regularization constant
- This penalty helps us avoid too complex candidates and hence reduce overfitting

Error of a model

- After training, the learning algorithm will return a model $h_o \in \mathcal{H}$
- How well does it work with future data? \Rightarrow Generalization!
- Maybe: $h_o \neq h_o^*$ and $h_o \neq h^*$
 - $h_o^* = \arg \min_{h \in \mathcal{H}} F(\mathcal{D}, h)$ is the *minimizer* of the empirical loss
 - $h^* = \arg \min_{h \in \mathcal{H}} F(P, h)$ is the *best member* of family \mathcal{H}
- Note:

$$F(P, h_o) - F(P, y^*) = F(P, h_o) - F(P, h^*) + F(P, h^*) - F(P, y^*)$$

- Estimation error: $F(P, h_o) - F(P, h^*)$
 - How good is the training algorithm?
- Approximation error: $F(P, h^*) - F(P, y^*)$
 - Capacity (representational power) of family \mathcal{H}



Error decomposition

- Estimation error

$$\begin{aligned} & F(P, h_o) - F(P, h^*) \\ &= F(P, h_o) - F(\mathbf{D}, h_o) + F(\mathbf{D}, h_o) - F(\mathbf{D}, h_o^*) + F(\mathbf{D}, h_o^*) - F(P, h^*) \end{aligned}$$

- It can be decomposed into two types of error: Optimization + Generalization
- Note that

$$|F(P, h_o) - F(\mathbf{D}, h_o)| \leq \sup_{h \in \mathcal{H}} |F(P, h) - F(\mathbf{D}, h)|$$

$$|F(\mathbf{D}, h_o^*) - F(P, h^*)| = \left| \min_{h \in \mathcal{H}} F(\mathbf{D}, h) - \min_{h \in \mathcal{H}} F(P, h) \right| \leq \sup_{h \in \mathcal{H}} |F(P, h) - F(\mathbf{D}, h)|$$

- Therefore

$$|F(P, h_o) - F(P, h^*)| \leq |F(\mathbf{D}, h_o) - F(\mathbf{D}, h_o^*)| + 2 \sup_{h \in \mathcal{H}} |F(P, h) - F(\mathbf{D}, h)|$$

(Optimization error)


(Generalization error)

Error decomposition

$$|F(P, h_o) - F(P, h^*)| \leq |F(\mathcal{D}, h_o) - F(\mathcal{D}, h_o^*)| + 2 \sup_{h \in \mathcal{H}} |F(P, h) - F(\mathcal{D}, h)|$$

- Estimation error can be decomposed into two types of error
- Optimization error: $F(\mathcal{D}, h_o) - F(\mathcal{D}, h_o^*)$
 - How close to optimality is h_o ?
 - h_o may not be the global minimum of the training loss
- Generalization error: $F(P, h_o) - F(\mathcal{D}, h_o)$
 - How far is the training loss from expected loss?
 - Smaller should be better
- In summary:

$Error(h_o) := \text{Optimization error} + \text{Generalization error} + \text{Approximation error}$






Errors by different factors

■ Data: manifold

- Complexity of the data space (độ phức tạp của không gian dữ liệu)

■ Data: distribution

- Complexity of the sampling distribution (độ phức tạp của phân bố lấy mẫu)
- Representativeness of the training set, ...




Errors by different factors (cont.)

■ Hypothesis space \mathcal{H}

- A bigger space (\mathcal{H}'), the (probably) smaller approximation error \rightarrow Why?
- More complex members, the (probably) smaller approximation error
 \rightarrow larger capacity
- An *effective* space (\mathcal{H}_e) is enough \rightarrow not too big/complex
- **Realizable** case: $F(P, h^*) = 0$

■ Training algorithm \mathcal{A}

- A better \mathcal{A} implies smaller estimation error
- A **bad** \mathcal{A} can provide small optimization error,
but large generalization error \rightarrow overfitting
- A good \mathcal{A} can localize an *effective* subset $\mathcal{H}^* \subset \mathcal{H}$



Errors by different factors (cont.)

- Each learner \mathcal{A} can have its own properties and biases
 - Some focus on *sparse solutions*, e.g., LASSO
 - Some can find simple models with *low complexity* (e.g., most weights ≈ 0)
 - Some often return models which are *robust w.r.t small changes in the input*
 - Some often return models which are *robust w.r.t weight noises*
 - Some are stable, ...
- Those biases should be used to analyze generalization
 - Generalization ability of a learning algorithm \mathcal{A}

Hypothesis space vs. Learner


- Consider the (unknown) regression function $y^*(\mathbf{x})$
- Let \mathcal{A}_D be the regressor, learned by method \mathcal{A} from a training set \mathbf{D}
- Note: We want that \mathcal{A}_D well approximates the truth y^* .
 - \mathcal{A}_D is random, according to the randomness when collecting \mathbf{D} .
- For any instance \mathbf{x} , the error made by \mathcal{A}_D is $(y^*(\mathbf{x}) - \mathcal{A}_D(\mathbf{x}))^2$
- ***The error made by the learner \mathcal{A} :***
(Lỗi của thuật toán \mathcal{A} khi phán đoán \mathbf{x})
$$err_{\mathcal{A}}(\mathbf{x}) = \mathbb{E}_{\mathbf{D}}(y^*(\mathbf{x}) - \mathcal{A}_D(\mathbf{x}))^2$$
- Why expectation? a different training set \mathbf{D}' will make \mathcal{A} to return a different function $\mathcal{A}_{D'}$,


The Bias-Variance Decomposition

$$err_A(x) = [Bias]^2 + Variance$$

- $Bias = y^*(x) - \mathbb{E}_{\mathcal{D}}[\mathcal{A}_{\mathcal{D}}(x)]; \quad Variance = \mathbb{E}_{\mathcal{D}} \left[(\mathcal{A}_{\mathcal{D}}(x) - \mathbb{E}_{\mathcal{D}'} \mathcal{A}_{\mathcal{D}'}(x))^2 \right]$
- This is known as **Bias-Variance Decomposition**

- ❖ Bias: how far is the **true value** from the **mean of predictions** by method \mathcal{A} ?
- ❖ Variance: how much does each prediction by \mathcal{A} vary around its mean?
- ❖ Small bias? Increase model complexity
→ Variance tends to increase
- ❖ Small variance? Decrease model complexity
→ Bias tends to increase






Error bounds

- Study theoretical guarantees by estimating the errors
(tạo ra các đảm bảo lý thuyết bằng việc ước lượng lỗi)
- **Approximation error:**

$$|F(P, y^*) - F(P, h^*)| \leq \epsilon_a$$

- Capacity of family \mathcal{H}
 - **Optimization error:**
- $$|F(\mathcal{D}, h_o^*) - F(\mathcal{D}, h_o)| \leq \epsilon_o$$
- Depending on the capacity of learning algorithm \mathcal{A}
 - and on the used optimizer



Bounds on Generalization Error

$$|F(P, h_o) - F(\mathcal{D}, h_o)| \leq \epsilon_g$$

- Generalizability of a learned function h_o
- Uniform bounds:


$$\sup_{h \in \mathcal{H}} |F(P, h) - F(\mathcal{D}, h)| \leq \epsilon_g$$

- Focus on the **worst** member
- May not be a good way to explain a learned function h_o

- PAC-Bayes bounds:

$$|\mathbb{E}_{h \in \mathcal{H}} [F(P, h) - F(\mathcal{D}, h)]| \leq \epsilon_g$$


- Study the error on average over \mathcal{H}
- May not explain a learned function h_o



Learnability

Generalization ability

- Consider a model $h_o \in \mathcal{H}$
 - Learned from a training set $\mathbf{D} \rightarrow h_o$ depends on \mathbf{D} ,
 - h_o is a function of \mathbf{D} , and we can write $h_o = h_o(\mathbf{D})$
- We want to estimate $F(P, h_o)$
 - *How well can h_o generalize on unseen data?*
 - **But P is unknown !!!**
 - Note that $F(P, h_o)$ depends on \mathbf{D}
 - When \mathbf{D} changes, so does $F(P, h_o)$
- Training set \mathbf{D}
 - Generated by a random process \rightarrow sometimes **good**, sometimes **bad** for training
 - There is always some probability that \mathbf{D} happens to be very *nonrepresentative* of the underlying distribution P



Probabilistic bounds

- An estimate of $F(P, h_o)$ should be in the form of **probabilistic bounds**

$$\Pr(F(P, h_o) > \epsilon) < \delta$$

- Interpret:

- The expected loss is greater than ϵ , with probability at most δ
(xác suất để lỗi trung bình vượt quá ϵ là nhỏ hơn δ)
- δ is the probability of getting a nonrepresentative sample \mathbf{D} , so that $F(P, h_o)$ is large
- ϵ is our prediction about the expected loss (phán đoán của ta về lỗi trung bình)

- Hope: *small ϵ and δ*

- Probability to sample a training set for which $F(P, h_o)$ is small

$$\Pr(F(P, h_o) \leq \epsilon) = \Pr(\{\mathbf{D}: F(P, h_o) \leq \epsilon\}) \geq 1 - \delta$$

- $1 - \delta$ represents **confidence** of our prediction ϵ



Approximately correct predictor

- When $F(P, h_o) \leq \epsilon$, we call h_o as *approximately correct predictor*
- The **failure** of the learner: $F(P, h_o) > \epsilon$
- We want to **upper bound** the probability of this event:

$$\Pr(F(P, h_o) > \epsilon)$$

- When \mathbf{D} contains m i.i.d samples from P , we can consider \mathbf{D} to be generated from distribution $P^m = P \times P \times \dots \times P$. Then

$$\Pr(F(P, h_o) > \epsilon) = P^m(\{\mathbf{D}: F(P, h_o) > \epsilon\})$$

- Tỉ lệ của vùng mà chứa các tập mẫu sẽ tạo ra lỗi lớn
- Sometimes we want to **lower bound**

$$\Pr(F(P, h_o) > \epsilon) > ???$$

Provably approximately correct (PAC)

■ **Learnability:** How to define it?

- PAC learning was introduced by Leslie Valiant (1984), the winner of the 2010 Turing Award for the introduction of the PAC model

■ **PAC learnability:** A hypothesis class \mathcal{H} is *PAC learnable* if there exist a learning alg. and a function $m_{\mathcal{H}}: (0,1) \times (0,1) \rightarrow \mathbb{N}$ with the following property:

For every $\epsilon, \delta \in (0,1)$, for every distribution P over \mathcal{X} , and for every labeling function $y^: \mathcal{X} \rightarrow \mathcal{Y}$, if the realizable assumption holds, then when running the learning algorithm on $m \geq m_{\mathcal{H}}(\epsilon, \delta)$ i.i.d. examples generated by P and labeled by y^* , the algorithm returns a hypothesis h_o such that $F(P, h_o) \leq \epsilon$ with probability of at least $1 - \delta$ (over the choice of the examples)*

$$\Pr(F(P, h_o) \leq \epsilon) \geq 1 - \delta$$

- Sau khi học từ một tập có ít nhất $m_{\mathcal{H}}$ mẫu, một thuật toán học sẽ trả về một hàm h_o thoả mãn $\Pr(F(P, h_o) \leq \epsilon) \geq 1 - \delta$
- *Realizable assumption:* $F(P, h^*) = 0$ for the best member $h^* = \arg \min_{h \in \mathcal{H}} F(P, h)$

Provably approximately correct (PAC)

$$\Pr(F(P, h_o) \leq \epsilon) \geq 1 - \delta$$

- The accuracy parameter ϵ determines how far h_o can be from the optimal one (this corresponds to the “approximately correct”),
- The confidence δ indicates how likely h_o is to meet that accuracy requirement (corresponds to the “probably” part of “PAC”).

Agnostic PAC learnability

- Realizable cases are rare in practice
 - They require \mathcal{H} to be too huge to “contain” the truth y^*
- **Non-realizable** cases: $F(P, h^*) \neq 0$ for the best member h^*
 - Hard problems
 - Small family \mathcal{H} , or each member is not complex enough
- **Agnostic PAC learnability:** *... when running the learning algorithm on $m \geq m_H(\epsilon, \delta)$ i.i.d. examples generated by P and labeled by y^* , the algorithm returns a hypothesis h_o such that, with probability at least $1 - \delta$ (over the choice of the examples):*

$$F(P, h_o) \leq F(P, h^*) + \epsilon$$

- No learner can learn a model with arbitrarily small error
- A learner can still declare success if its error is not much larger than the best error achievable by a predictor from \mathcal{H}

Sample complexity

- **Sample complexity:** the minimal number of examples are required to guarantee a probably approximately correct solution
- **Lemma:** if the realizable assumption holds, every finite hypothesis class \mathcal{H} is PAC learnable with sample complexity $m_H(\epsilon, \delta) \leq \lceil \frac{\log(|\mathcal{H}|/\delta)}{\epsilon} \rceil$
- $m_H(\epsilon, \delta)$ is often required to be *polynomial* in $1/\epsilon$ and $\log(\delta)$
 - to say learnability of \mathcal{H} [Valiant, 1984]
- **The curse of dimensionality:** $m_H(\epsilon, \delta)$ is *exponential* in the dimensionality n
(lời nguyền của số chiều: nếu lượng mẫu cần thiết có cỡ cấp hàm mũ của n)
 - \mathcal{H} không thực tế!

No-free-lunch theorem

- For one problem, there are many learners
 - Some are really powerful, e.g., using SGD
 - More training samples can help training to be loss closer and closer to expected loss
 - the trained model can be better
- **Is there a universal learner that can learn any functions** just by increasing training size?
- **No-free-lunch theorem:** *Let \mathcal{A} be any learner for the task of binary classification with input space \mathcal{X} , the 0-1 loss, and any training size $m \leq |\mathcal{X}|/2$. Then there exists a distribution P over $\mathcal{X} \times \{0, 1\}$ satisfying:*
 - *There exists a function h with $F(P, h) = 0$*
 - *$F(P, \mathcal{A}(\mathbf{D})) \geq 1/8$, with probability of at least $1/7$ over the choice of $\mathbf{D} \sim P^m$*
- **No universal learner !**

Assessing the Generalization ability



Basic bounds

Generalization ability

- Consider a model $h_o \in \mathcal{H}$
 - Learned from a training set \mathbf{D}
 - h_o is a function of \mathbf{D} , and we can write $h_o = h_o(\mathbf{D})$
- We want to estimate $F(P, h_o)$
 - *How well can h_o generalize on unseen data?*
 - **But P is unknown !!!**
 - Need to assure

$$\Pr(F(P, h_o) \leq F(P, h^*) + \epsilon) \geq 1 - \delta$$
- Another way: estimate **Generalization error**

$$F(P, h_o) - F(\mathbf{D}, h_o)$$



Bounds on Generalization error

$$F(P, h_o) - F(\mathbf{D}, h_o)$$

- $F(\mathbf{D}, h_o)$ is computed from \mathbf{D} with m samples, and is an empirical version of $F(P, h_o)$
- By the law of large numbers:
$$\Pr\left(\lim_{m \rightarrow \infty} F(\mathbf{D}, h_o) = F(P, h_o)\right) = 1$$
 - But this law does not provide an explicit form of the *convergence rate*
- **Theorem (Hoeffding):** *Let z_1, z_2, \dots, z_m be i.i.d. random variables and function $g(z) \in [a, b]$. Then for all $\epsilon > 0$, we have*

$$\Pr\left(\mathbb{E}_z[g(z)] - \frac{1}{m} \sum_{i=1}^m g(z_i) > \epsilon\right) \leq \exp\left(-\frac{2m\epsilon^2}{(b-a)^2}\right)$$

- Denote the right hand side as δ , then

$$\Pr\left(\mathbb{E}_z[g(z)] - \frac{1}{m} \sum_{i=1}^m g(z_i) > (b-a) \sqrt{\frac{1}{2m} \log \frac{1}{\delta}}\right) \leq \delta$$

Basic bounds: Hoeffding's inequality

- Or with probability at least $1 - \delta$, we have

$$\mathbb{E}_z[g(z)] - \frac{1}{m} \sum_{i=1}^m g(z_i) \leq (b - a) \sqrt{\frac{1}{2m} \log \frac{1}{\delta}} \quad (1)$$

- This is true for any *fixed* function g (đúng đối với một hàm g cố định)
- However, it does not apply when g depends on the sampling of the data (không đúng nếu g phụ thuộc vào tập z_1, z_2, \dots, z_m)
- **Limitation:** for a given fixed function g
 - There exists a sample \mathbf{D} of size m , so that inequality (1) holds
 - The set S of those samples has measure $\Pr(\mathbf{D} \in S) \geq 1 - \delta$
 - However, a different function g' may have a different S'
 - For a given sample \mathbf{D} , only some members of \mathcal{H} can satisfy (1)
→ a model h_o trained from \mathbf{D} may not satisfy (1)

Uniform bound

- Before seeing the data, which function will the algorithm choose?
 - We do not know!
- Consider

$$\sup_{h \in \mathcal{H}} [F(P, h) - F(\mathcal{D}, h)] \quad (2)$$

- Note: $F(P, h_o) - F(\mathcal{D}, h_o) \leq \sup_{h \in \mathcal{H}} [F(P, h) - F(\mathcal{D}, h)]$
- If we can upper bound (2), we are done!
 - This bound on error will hold true for every member of \mathcal{H}
(cận trên này về lỗi sẽ đúng đối với mọi hàm trong \mathcal{H})
 - It tries to estimate the error of the **worst** member !!!
(tức là ta xem xét lỗi của thành viên tệ nhất)

Uniform bound for finite hypothesis class

- Assume that \mathcal{H} is finite, and $|\mathcal{H}| = N$ (tập \mathcal{H} có N phần tử)
- For each $h_i \in \mathcal{H}$, let \mathbf{D}_i be a ‘bad’ sample so that $F(P, h_i) - F(\mathbf{D}_i, h_i) > \epsilon$
 - Hoeffding’s inequality: $\Pr(\mathbf{D}_i) = \Pr(\mathbf{D}_i : F(P, h_i) - F(\mathbf{D}_i, h_i) > \epsilon) \leq \exp\left(-\frac{m\epsilon^2}{2C^2}\right)$
(C is the maximum value of the loss function)
- Union bound: $\Pr(\mathbf{D}_i \cup \mathbf{D}_j) \leq \Pr(\mathbf{D}_i) + \Pr(\mathbf{D}_j)$
- Generally: $\Pr(\mathbf{D}_1 \cup \mathbf{D}_2 \cup \dots \cup \mathbf{D}_N) \leq \sum_{i=1}^N \Pr(\mathbf{D}_i)$
- As a result: $\Pr(\exists h \in \mathcal{H} : F(P, h) - F(\mathbf{D}, h) > \epsilon) \leq \sum_{i=1}^N \Pr(\mathbf{D}_i) \leq N \exp\left(-\frac{m\epsilon^2}{2C^2}\right)$
- Hence, for all $\delta > 0$, with probability at least $1 - \delta$:

$$\forall h \in \mathcal{H}, \quad F(P, h) \leq F(\mathbf{D}, h) + C \sqrt{\frac{2 \log N - 2 \log \delta}{m}}$$

PAC learnability: a revisit

- Consider realizable cases and finite hypothesis class \mathcal{H}
- Let ERM as the learner that always returns $h_{erm}^* = \arg \min_{h \in \mathcal{H}} F(\mathbf{D}, h)$
 - The realizable assumption implies $F(\mathbf{D}, h_{erm}^*) = 0$
- Let \mathcal{H}_{erm} contain all functions returns by this learner
 - Then, for all $\delta > 0$, with probability at least $1 - \delta$:

$$\forall h_{erm} \in \mathcal{H}_{erm}, \quad F(P, h_{erm}) \leq C \sqrt{\frac{2 \log N - 2 \log \delta}{m}}$$

- For all $\epsilon > 0$, letting $m_H = C^2(2 \log N - 2 \log \delta)/\epsilon^2$, for every $m \geq m_H$:

$$\Pr(F(P, h_{erm}) \leq \epsilon) \geq 1 - \delta$$

- \mathcal{H} is PAC learnable!

$$\forall h \in \mathcal{H}, \quad F(P, h) \leq F(\mathcal{D}, h) + C \sqrt{\frac{2 \log N - 2 \log \delta}{m}}$$

This bound does not apply to the infinite case

$$N = |\mathcal{H}| = \infty$$

Different approaches

- Performance of a model depends on:
 - The learning algorithm \mathcal{A}
 - The model family \mathcal{H}
 - Data distribution P
 - Labeling function y^*
- Approaches to bounding generalization ability:
 - Complexity/capacity of family \mathcal{H} \rightarrow Complexity-based bounds
 - Capacity of the learning algorithm \mathcal{A} \rightarrow Algorithmic-based bounds
 - Complexity of the sampling distribution P \rightarrow Data-dependent bounds
 - Capacity of the trained model \rightarrow Model-dependent bounds

Algorithmic robustness

- Robust learner can receive a training set \mathbf{D} to
 - produce a model which has similar loss around the training samples
 - return a model which can generalize well to the local areas around \mathbf{D}
- [Xu & Mannor, 2012] formalize the concept of Robustness
- **Definition:** A learning algorithm \mathcal{A} is (K, ϵ) -robust, for $K \in \mathbb{N}$ and $\epsilon: \mathcal{Z}^m \rightarrow \mathbb{R}$, if \mathcal{Z} can be partitioned into K disjoint sets, denoted by $\{\mathcal{Z}_k\}_{k=1}^K$, such that the following holds for all $\mathbf{D} \in \mathcal{Z}^m$:
$$\forall \mathbf{s} \in \mathbf{D}, \forall \mathbf{z} \in \mathcal{Z}, \text{if } \mathbf{s}, \mathbf{z} \in \mathcal{Z}_k \text{ for some } k \in [K] \text{ then } |\ell(\mathcal{A}_D, \mathbf{s}) - \ell(\mathcal{A}_D, \mathbf{z})| \leq \epsilon(\mathbf{D})$$
 - Where $\ell(\mathcal{A}_D, \mathbf{s})$ is the loss of model \mathcal{A}_D at example \mathbf{s} , and $\mathcal{Z} = \mathcal{X} \times \mathcal{Y}$ for supervised problems
 - Every model returned from \mathcal{A} should be robust at some small areas around the training points (mỗi mô hình học được sẽ vững trước sự thay đổi nhỏ xung quanh các mẫu học)
 - A *non-robust* model will be sensitive to some small changes in the input
→ Vulnerable to Adversarial Attacks

Robustness-based bound

- **Theorem:** Assume that \mathcal{A} is (K, ϵ) -robust. Let \mathbf{D} contain m i.i.d samples from distribution P , and h is trained from \mathbf{D} by \mathcal{A} , and f be the loss with max value C . For any $\delta > 0$, with probability at least $1 - \delta$,

$$F(P, h) \leq F(\mathbf{D}, h) + \epsilon(\mathbf{D}) + C \sqrt{\frac{2K \log 2 + 2\log(1/\delta)}{m}} \quad (9)$$

- When both ϵ and training loss is small, the expected loss can be small
→ h can generalize well and avoid adversarial attacks
- A robust learner can return models with high generalization ability
but the reverse is not true!
- *Tradeoff:* increasing K can reduce ϵ
- Idea to improve generalization
- Penalize the robustness level ϵ when training h → **Robust training, adversarial training**

Robust algorithms

- **Learners with Lipschitz loss:** if the instance set \mathcal{Z} is compact and the loss f is L -Lipschitz continuous in the input, then \mathcal{A} is $(\mathcal{N}(\gamma/2, \mathcal{Z}), L\gamma)$ -robust for all $\gamma > 0$.
 - $\mathcal{N}(\gamma/2, \mathcal{Z})$ is the covering number, i.e., the minimal number of balls with radius at most $\gamma/2$ that cover \mathcal{Z} completely
(số lượng ít nhất các hình cầu có bán kính không quá $\gamma/2$ mà có thể bao toàn bộ \mathcal{Z})
- **LASSO:** $(\mathcal{N}(\gamma/2, \mathcal{Z}), \gamma \frac{1}{m} \sum_{i=1}^m y_i^2 / \lambda + \gamma)$ -robust for all $\gamma > 0$
- **PCA:** $(\mathcal{N}(\gamma/2, \mathcal{Z}), 2\gamma dB)$ -robust, if we use d components and B is the diameter of \mathcal{Z}
- **Max margin learners:** $(\mathcal{N}(\gamma/2, \mathcal{Z}), \gamma)$ -robust
 - where γ is the **margin** of \mathcal{A}_D :

$$\mathcal{A}_D(x) = \mathcal{A}_D(s) \quad \text{for any } x \in \mathcal{X} \text{ and } s \in \mathcal{D} \text{ in the same ball with diameter } \gamma$$
- **Learners for neural networks:** $(\mathcal{N}(\gamma/2, \mathcal{Z}), \gamma W^d)$ -robust
 - For DNNs with d layers, each weight matrix has norm at most W

$$\begin{aligned} & \text{Minimize} \\ & F(\mathbf{D}, h) + \lambda \|\mathbf{w}\|_1 \end{aligned}$$

Algorithmic robustness: issues

$$F(P, h) \leq F(\mathcal{D}, h) + \epsilon(\mathcal{D}) + C \sqrt{\frac{2K \log 2 + 2\log(1/\delta)}{m}}$$

- K is often large to ensure a small ϵ
 - Previous examples: $K = \mathcal{N}(\gamma/2, \mathcal{Z})$ → exponential in the dimension
 - \rightarrow Vacuous/trivial bound
- Vacuousness also come from ϵ
 - By definition, $\epsilon(\mathcal{D}) = \sup\{|\ell(h, \mathbf{s}) - \ell(h, \mathbf{z})| : \mathbf{s} \in \mathcal{D}, \mathbf{z} \in \mathcal{Z}_k\}$ where \mathcal{Z}_k contains \mathbf{s}
 - For 0-1 loss, one incorrect prediction may lead to $\epsilon = 1$ → Vacuous