

# Non-linear dimensionality reduction and kernels: eigenmaps, isomaps, locally linear embeddings

# Review of Dimensionality Reduction

- Dimensionality reduction can be done through
  1. **feature selection**: only keeps the most relevant variables from the original dataset.
  2. **dimensionality reduction**: finds the smaller set of new variables, containing basically the same information as the original variables.
- Dimensionality reduction can also be categorized into:
  - linear dimensionality reduction (e.g. PCA, SVD)
  - non-linear dimensionality reduction (e.g. autoencoders, kernel PCA and others).

# Non-linear dimensionality reduction

- **This lecture:** find a nonlinear low dimensional representation of data that reflects the data topology
- **Methods covered:**
  - Isomaps
  - Locally Embedding Space (LLE)
  - Eigenmaps

# Motivation

While the data is represented in high dimensions, it usually only has a few degrees freedom.

Examples:

- Scanned images of handwritten characters

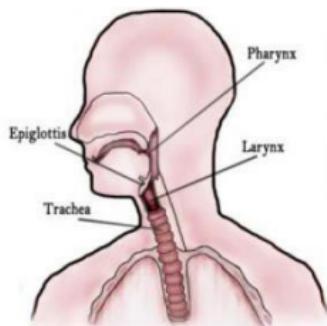


Representation:  $28 \times 28$  pixels (= 784 dim)

Few attributes such as: shape, tilt and cursiveness govern the actual written character.

- Natural processes with physical constraints - speech

Few anatomical characteristics, such as size of the vocal chords, pressure applied, etc. govern the speech signal.



# Topology

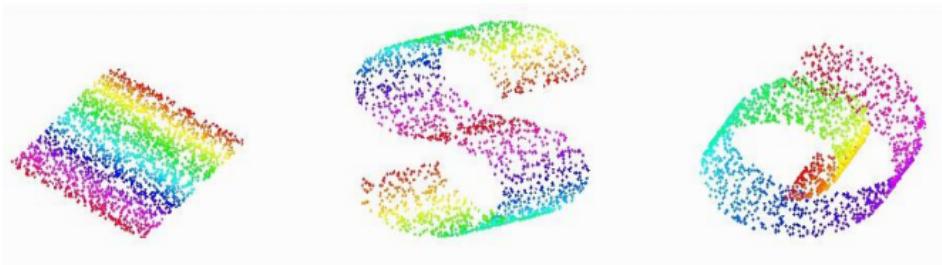
- Topology: “*the study of qualitative properties of certain objects that are invariant under a certain kind of transformation, especially those properties that are invariant under a certain kind of invertible transformation*”
- “A topologist is one who doesn’t know the difference between a doughnut and a coffee cup” –John L. Kelley (In General Topology 1995, 88 footnote)

# Manifolds

- Definition: A manifold is a topological space that locally resembles Euclidean space **near each point.**



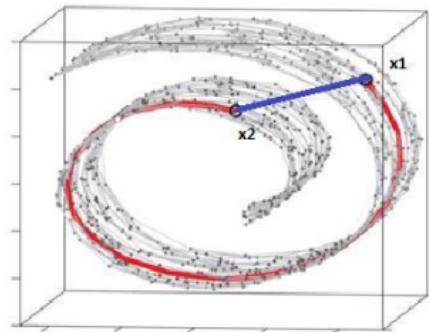
# Manifolds



- Three examples of manifolds
- All three are two-dimensional data embedded in 3D
  - Linear, “S”-shape, “Swiss roll”
- For all three examples, we would like to recover:
  - Their two-dimensional representation
  - “Consistent” coordinates of the data in the 2D

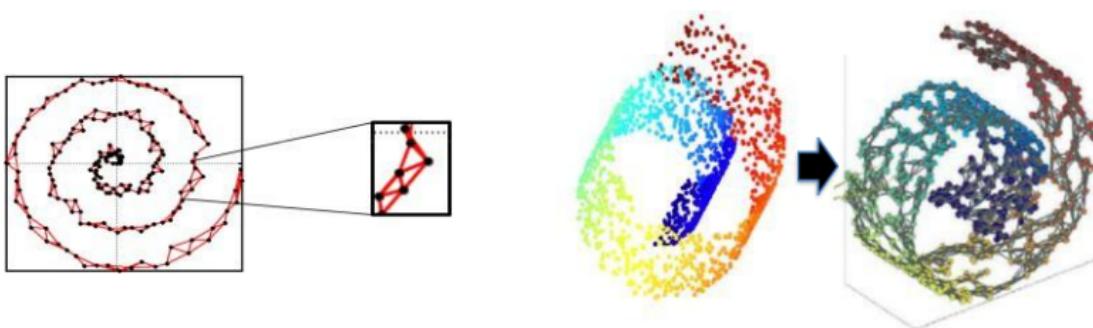
# Manifolds and local distances

- In general, the distances induced by data (manifolds) may not be Euclidean. That is the global distances don't respect the geometry.
- Local distances can be still approximated with Euclidean distances
- **Idea for the dimensionality reduction:**
  - Define global distances/similarity in terms of local distances/similarity
  - Use these to define a low-dimensional embedding (low-dimensional representation) of the data
- **The idea can be implemented with the help of the Neighborhood graph**



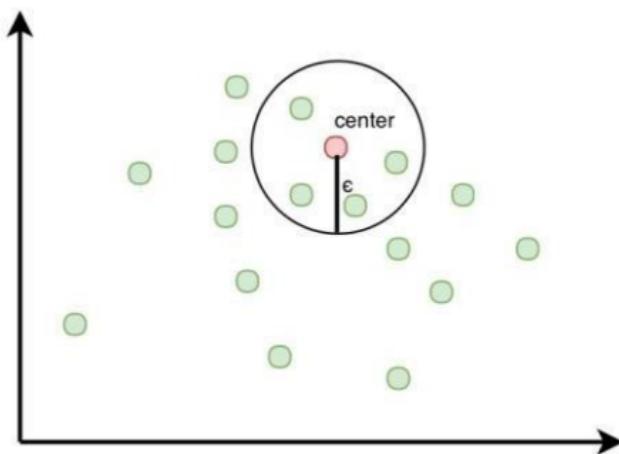
# Neighborhood Graph

- A neighborhood graph
  - Vertices = data points
  - Edges and their weights reflect local similarity or local distances
  - Only points close to each other (neighbors) are connected



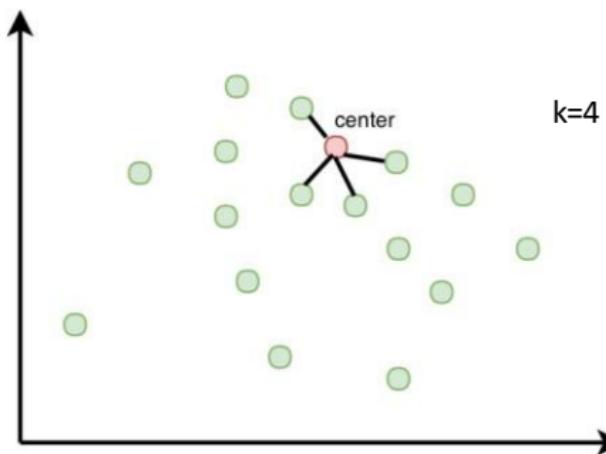
# Neighborhood Graph

**Approach 1:** select and connect all points within the  $\epsilon$  neighborhood



# Neighborhood Graph construction

**Approach 2:** pick and connect k nearest neighbor points



# Neighborhood Graph

- **Distance-based neighborhood graph:**
  - Edge weight: Euclidean distance between two data points –  $d(x_i, x_j)$
- **Similarity-based neighborhood graph:**
  - Edge weight:
    - Simple:  $W_{ij} = 1$  if connected, 0 otherwise
    - Kernel:  $W_{i,j} = \exp\left(-\frac{\|x_i - x_j\|^2}{t}\right)$  if connected, 0 otherwise

# Non-linear dimensionality reduction

- Three methods to define lower dimensional embedding of data instances:
  - Isomaps
  - Locally Embedding Space (LLE)
  - Eigenmaps
- All of these rely on the neighborhood graph connecting only data instances close to each other
- First let us introduce Multi-dimensional scaling (MDS) method ...

# Multidimensional Scaling (MDS)

- MDS is a classical approach that can map the original high dimensional space to a lower dimensional space. It attempts to preserve the pairwise distance among the data points.
- It is used when we want to visualize high dimensional data say in 2 or 3D

# Multidimensional Scaling (MDS)

**Idea:** MDS maps points to a low dimensional space (say of dimension  $k$ ) such that the Euclidean distances between the points in this new space approximate the original distance matrix.

$$\Delta := \begin{pmatrix} \delta_{1,1} & \delta_{1,2} & \dots & \delta_{1,I} \\ \delta_{2,1} & \delta_{2,2} & \dots & \delta_{2,I} \\ \vdots & \vdots & & \vdots \\ \delta_{I,1} & \delta_{I,2} & \dots & \delta_{I,I} \end{pmatrix}$$

Map input points  $x_i$  to  $z_i$  such that

- Classical MDS: the norm  $\|z_i - z_j\|$  is the Euclidean distance
- Objective function:  $\|\cdot\|$

$$\min_{z_1, \dots, z_I} \sum_{i < j} (\|z_i - z_j\| - \delta_{i,j})^2$$

# Isomap

## Algorithm

- **Step 1:** Construct the neighborhood graph  $G$  with edge weights corresponding to local distances
- **Step 2:** compute the shortest distance between all pairs of points:
  - The shortest distance can be calculated via Floyd or Dijkstra algorithm.

$$\Delta := \begin{pmatrix} \delta_{1,1} & \delta_{1,2} & \dots & \delta_{1,J} \\ \delta_{2,1} & \delta_{2,2} & \dots & \delta_{2,J} \\ \vdots & \vdots & & \vdots \\ \delta_{I,1} & \delta_{I,2} & \dots & \delta_{I,J} \end{pmatrix}$$

- **Step 3:** Construct the  $k$ -dimensional embedding
  - Use classical MDS to find a  $k$ -dimensional embedding

$$\min_{z_1, \dots, z_I} \sum_{i < j} (\|z_i - z_j\| - \delta_{i,j})^2$$

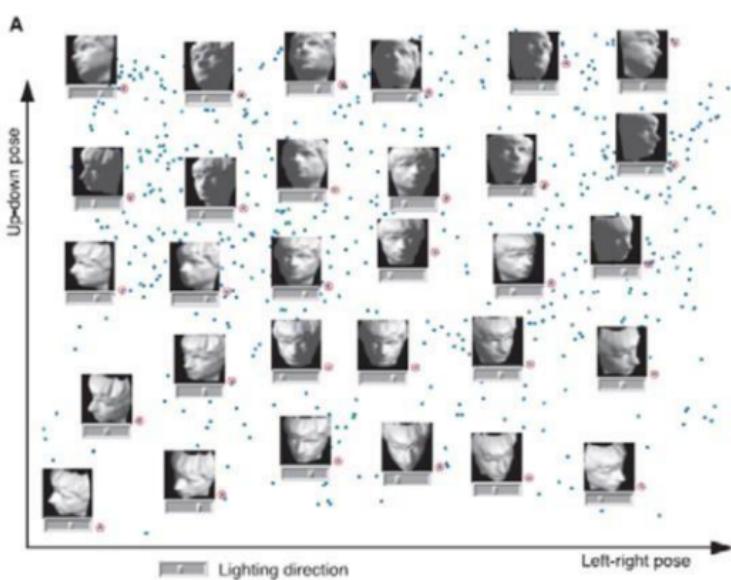
All  $z_1, z_2, \dots, z_I$  are  $k$ -dimensional

# Isomap

- **Advantages:**
  - Non-linear dimensionality reduction
  - Non-iterative polynomial time algorithm
  - Guarantee of globally optimality:
    - For intrinsically Euclidean manifolds, a guarantee of asymptotic convergence to the true structure
    - The ability to discover manifolds of arbitrary dimensionality
- **Disadvantage:**
  - Sensitive to noise
  - Few free parameters

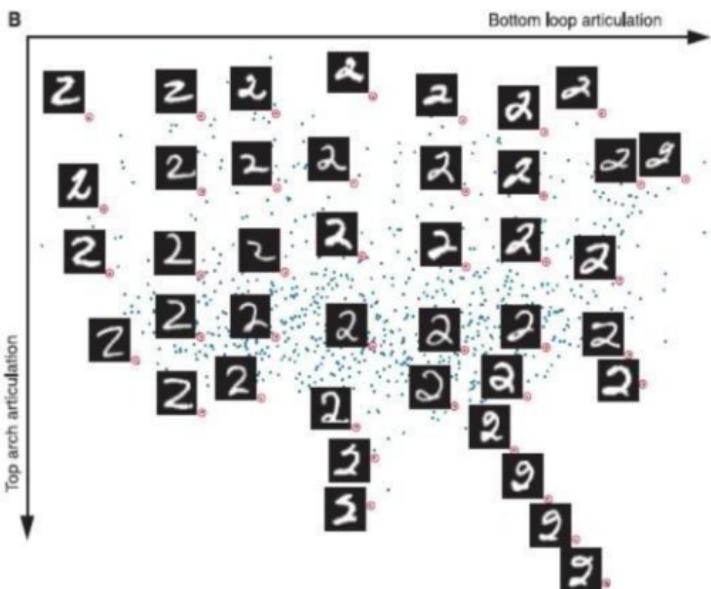
# Isomaps: Example

- Dimensionality Reduction for visual perception
  - 64 x 64 image
  - 698 raw images
  - Isomap ( $k = 6$ )



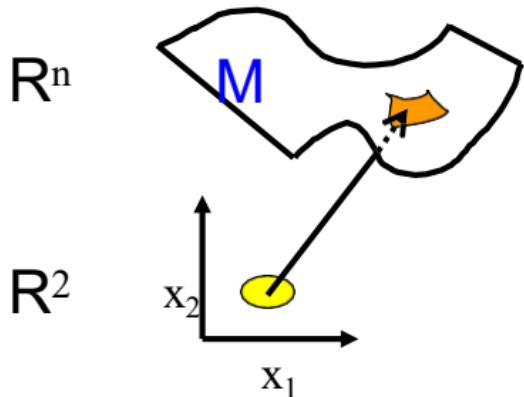
# Isomap: Example

- Handwritten '2'
  - 1000 handwritten 2s
  - Isomap ( $\epsilon = 4.2$ )



# Locally Linear Embedding (LLE)

- Manifold Characteristics/Key Assumption
  - We expect each data point and its neighbors to lie on or close to a **locally linear patch**
  - But, how to combine all local patches together?

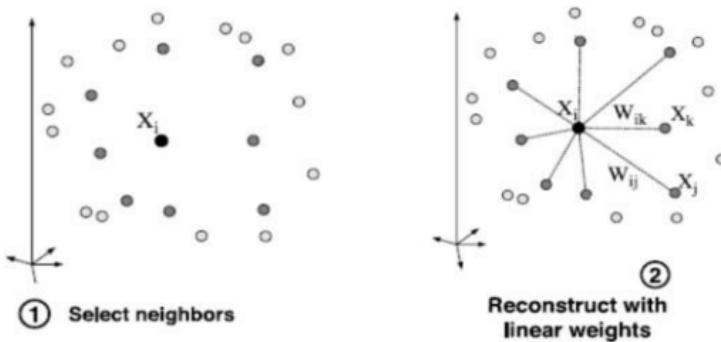


# LLE: Intuition

- Assume that manifold is approximately “linear” when viewed locally (in a small neighborhood)
- A good projection should preserve this local geometric property as much as possible

# LLE algorithm

- **Step 1:** Select neighbors for each data instance  $x_i$
- **Step 2:** Each data instance is written as a convex combination of its neighbors. Weights of the convex combination ‘reconstruct’ each point from its neighbors.



# LLE Algorithm

- **Step 2:** The weights chosen aim to minimize the reconstruction error.

$$\min_W \left\| X_i - \sum_{j=1}^K W_{ij} X_j \right\|^2$$

Note: Assign weights under two constraints:

- $W_{ij} = 0$  if  $X_j$  does not belong to set of neighbors of  $X_i$
- The rows of the weight matrix sum to one i.e.

$$\sum_j W_{ij} = 1$$

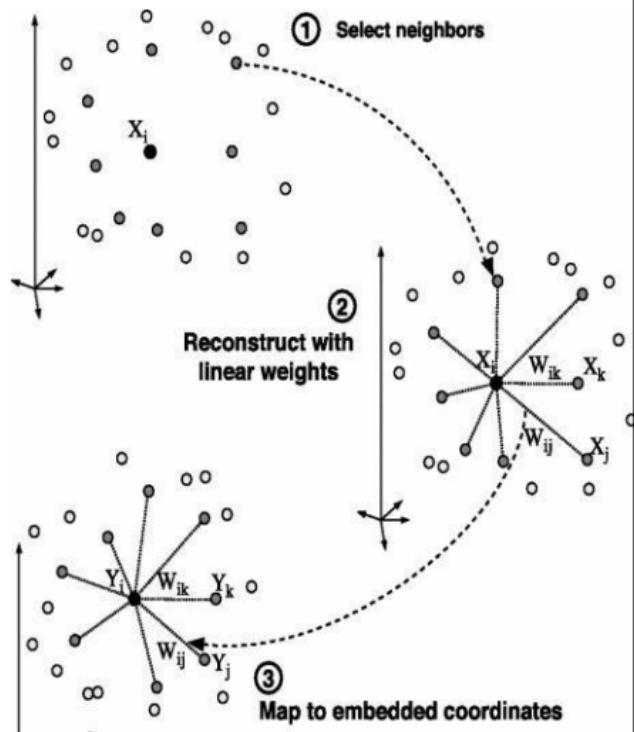
# LLE Algorithm

- **Step 3:** Map  $R^n$  to a low-dimensional embedding  $R^k$

$$\begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{pmatrix}^{R^n} \rightarrow \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{pmatrix}^{R^k}$$

The cost function can be minimized by solving a sparse NxN eigenvalue problem:

$$\min_y \sum_{i=1} \left\| Y_i - \sum_j W_{ij} Y_j \right\|^2$$



# LLE: Eigenvalue Problem

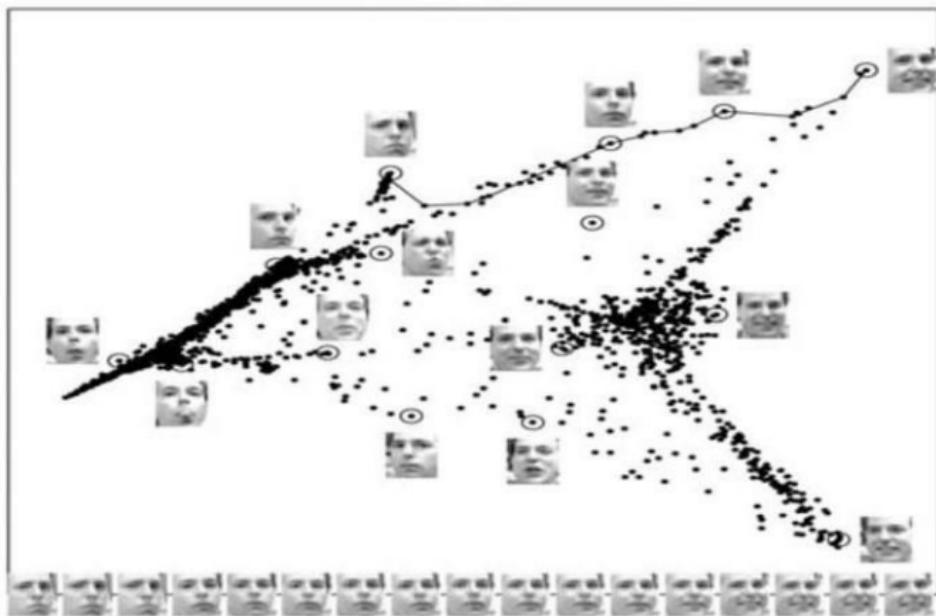
The following is a more direct and simpler derivation for Y:

Define  $M = (I - W)^T (I - W)$

$$\begin{aligned}\Phi(Y) &= \sum_i \left\| Y_i - \sum_j W_{ij} Y_j \right\|^2 = \sum_i \left\| Y_i - [Y_1, Y_2, \dots, Y_n] W_i^T \right\|^2 \\ &= \left\| [Y_1, Y_2, \dots, Y_n] - [Y_1, Y_2, \dots, Y_n] [W_1^T, W_2^T, \dots, W_n^T]^T \right\|_F^2 \\ &= \left\| Y - YW^T \right\|_F^2 = \left\| Y(I - W^T) \right\|_F^2 = \text{trace}(Y(I - W)^T (I - W) Y^T) \\ &= \text{trace}(YMY^T)\end{aligned}$$

where  $Y = [Y_1, Y_2, \dots, Y_n]$

# LLE Example



Images of faces mapped into the embedding space described by the first two coordinates of LLE. Representative faces are shown next to circled points. The bottom images correspond to points along the top-right path (linked by solid line) illustrating one particular mode of variability in pose and expression.

# LLE: effect of k

- Require dense data points on the manifold for good estimation

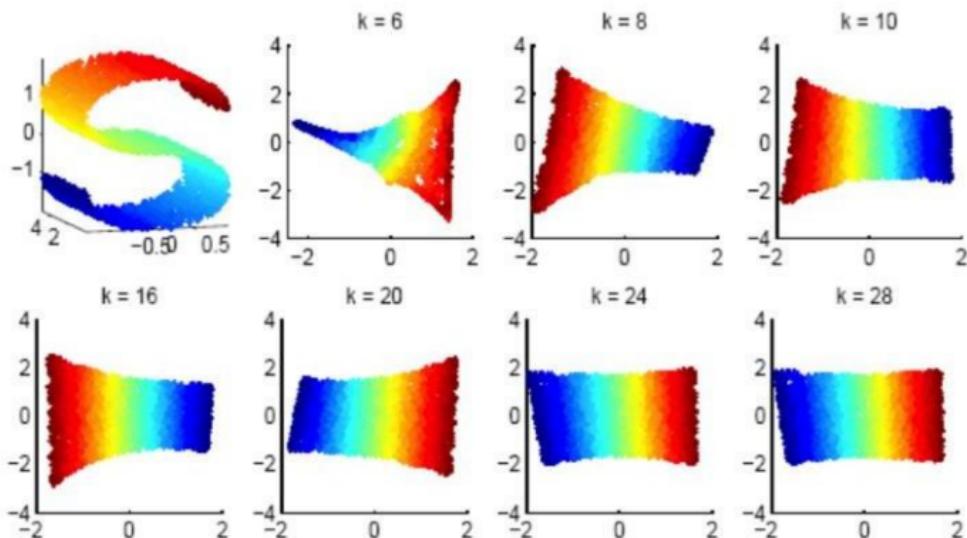


FIG. 5. S-curve (top left) and computed 2D coordinates by LLE with various neighborhood size  $k$ .

# Limitations of LLE

- require dense data points on the manifold for good estimation
- A good neighborhood seems essential to their success
  - How to choose k?
  - Too few neighbors
    - Results in rank deficient tangent space and lead to over-fitting
  - Too many neighbors
    - Tangent space will not match local geometry well

# Laplacian Eigenmaps

- Problem: Given a set  $(x_1, x_2, \dots, x_n)$  of  $n$  points in  $\mathbb{R}^d$ , find a set of points  $(y_1, y_2, \dots, y_n)$  in  $\mathbb{R}^k$  ( $k \ll d$ ) such that  $y_i$  represents  $x_i$ .
- Steps
  - Build the adjacency graph
  - Choose the weights for edges in the graph
  - Eigen-decomposition of the graph Laplacian
  - Form the low-dimensional embedding

# Laplacian Eigenmaps Algorithm

- **Step 1:** Construct the neighborhood graph with weights modeling local similarities:
  - Simple: 1 if connected; 0 otherwise
  - Kernels: Gaussian or Heat kernels

$$W_{ij} = e^{-\frac{\|x_i - x_j\|^2}{t}}$$

- **Step 2:** Construct Graph-Laplacian matrix
  - Construct diagonal weight matrix D from the weight matrix
  - Construct Laplacian matrix L = D-W
    - Laplacian is a symmetric, positive semi-definite matrix

# Laplacian Eigenmaps Algorithm

- **Step 3:** Finding of the lower dimensional embedding
  - Find  $Y = (y_1, y_2, \dots, y_n)^T$  that preserves local similarities
  - Note: each  $y$  is low-dimensional

**Objective:** minimize

$$\sum_{ij} (y_i - y_j)^2 W_{ij}$$

- The above objective function can be rewritten as:

$$\frac{1}{2} \sum_{i,j} (y_i - y_j)^2 W_{ij} = Y^T LY$$

- **Solution:** Compute eigenvalues and eigenvectors of the generalized eigenvector problem

# Generalized Eigenvector Problem

- Suppose we have the  $n$  points such that  $(x_1, x_2, \dots, x_n)$  in  $\mathbb{R}^d$
- Construct the weight matrix between each point

$$W = \begin{pmatrix} w_{11} & w_{12} & \cdots & w_{1n} \\ w_{21} & w_{22} & \cdots & w_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ w_{n1} & w_{n1} & \cdots & w_{nn} \end{pmatrix}$$

- Based on equation  $|W - \lambda I| = 0$ , we can compute the Eigenvalues  $\lambda = (\lambda_0, \lambda_1, \dots, \lambda_n)$

# Euclidean Embedding Space

Once we have all the Eigenvalues, then we solve equation  $Lv = \lambda Dv$  to compute the Eigenvector  $v$

$$Lv = \lambda Dv$$

$$Lv - \lambda Dv = 0$$

$$(L - \lambda D)v = 0$$

Since  $(L - \lambda D)$  has the dimension  $n * n$

Then Eigenvector  $v$  has the dimension  $n * 1$

# Euclidean Embedding Space

- Let  $v_0, v_1, \dots, v_{n-1}$  be the eigenvectors solutions to the equation  $Lv = \lambda Dv$ , ordered according to their eigenvalues,

$$Lv_0 = \lambda_0 Dv_0$$

$$Lv_1 = \lambda_1 Dv_1$$

...

$$Lv_{n-1} = \lambda_{n-1} Dv_{n-1}$$

$$Lv_n = \lambda_n Dv_n$$

$$0 = \lambda_0 \leq \lambda_1 \leq \dots \leq \lambda_{n-1}$$

# Euclidean Embedding Space

- Then we leave out the eigenvector  $v_0$  corresponding to eigenvalue  $\lambda_0$  and use the next  $k$  eigenvectors  $(v_1, v_2, \dots, v_k)$  for embedding in  $k$ -dimensional Euclidean space.
- The  $k$ -dimensional embedding space is denoted as  $Y = (y_1, y_2, \dots, y_k)^T$

# Euclidean Embedding Space

Once we find all Eigenvectors, then, we can project the data points to lower dimension  $R^k$  embedding space

$$X = (x_1, x_2, \dots, x_n)^{R^d} \rightarrow Y = \begin{pmatrix} y_1 & y_2 & \vdots & y_k \\ v_1(1) & v_2(1) & \vdots & v_k(1) \\ v_1(2) & v_2(2) & \ddots & v_k(2) \\ \vdots & \vdots & \ddots & \vdots \\ v_1(n) & v_2(n) & \dots & v_k(n) \end{pmatrix}^T$$

We reduce the dimension from  $n * d$  to  $n * k$

$y_1 = (v_1(1), v_1(2), \dots, v_1(n))$  represents the coordinates of all  $n$  data points at 1<sup>st</sup> dimension in the embedding space.

# Optimal embedding space

- Objective function:

$$\min \frac{1}{2} \sum_{i,j} (y_i - y_j)^2 W_{ij} \quad \underset{\substack{y^T D y = 1 \\ y^T D 1 = 0}}{\operatorname{argmin}} y^T L y$$

- The constraint  $y^T D y = 1$  removes arbitrary scaling factor
- The constraint  $y^T D 1 = 0$  removes a translation invariance in  $y$

# Laplacian Eigenmaps: Embedding Space

After the optimization, each data point can be mapped into the optimal k-dimensional embedding space

$$Y = (y_1, y_2, \dots, y_k)^T$$

$$x_i^{R^{dn}} \rightarrow (v_1(i), v_2(i), \dots, v_k(i))$$

$v_j(i)$  is the coordinate of point  $x_i$  at jth dimension at k-dimensional space, where  $1 \leq j \leq k$

# Laplacian Eigenmaps Example

- Suppose we want to project  $X = \{x_1, x_2, \dots, x_n\}^{R^m}$  into 2 dimensional space
- Let  $\{\lambda_0, \lambda_1, \lambda_2, \dots, \lambda_n\}$  to be the calculated eigenvalues with increasing order  $0 = \lambda_0 \leq \lambda_1 \leq \dots \leq \lambda_n$
- Let  $v = \{v_1, v_2, \dots, v_n\}$  be the eigenvector solutions of equation  $Lv = \lambda Dv$
- Take  $\{v_1, v_2\}$  for embedding in 2-dimensional space

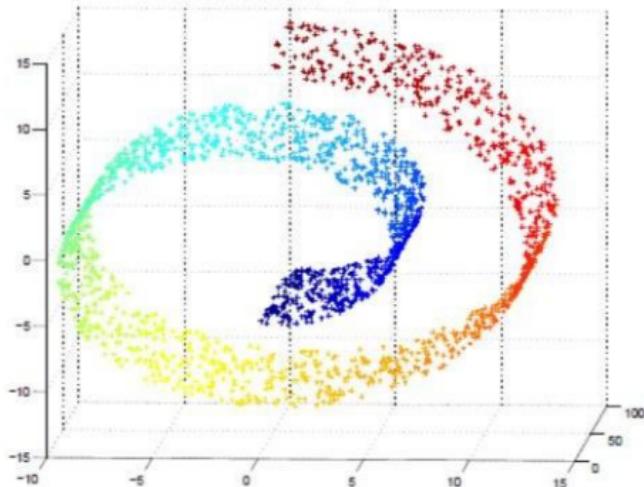
$$\begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{pmatrix}^m \rightarrow \begin{pmatrix} v_1(1) & v_2(1) \\ v_1(2) & v_2(2) \\ \vdots & \vdots \\ v_1(n) & v_2(n) \end{pmatrix}^2$$

# Locally Linear Embedding and Laplacian Eigenmap

- LLE is connected with Laplacian Eigenmap
- LLE minimizes  $Y^T(I-W)^T(I-W)y$  to reduce eigenvectors of  $(I-W)^T(I-W)$
- Actually, finding eigenvectors of  $(I-W)^T(I-W)$  can be re-interpreted as finding eigenvectors of iterated Laplacian  $L^2$ .

# Laplacian Eigenmap Example

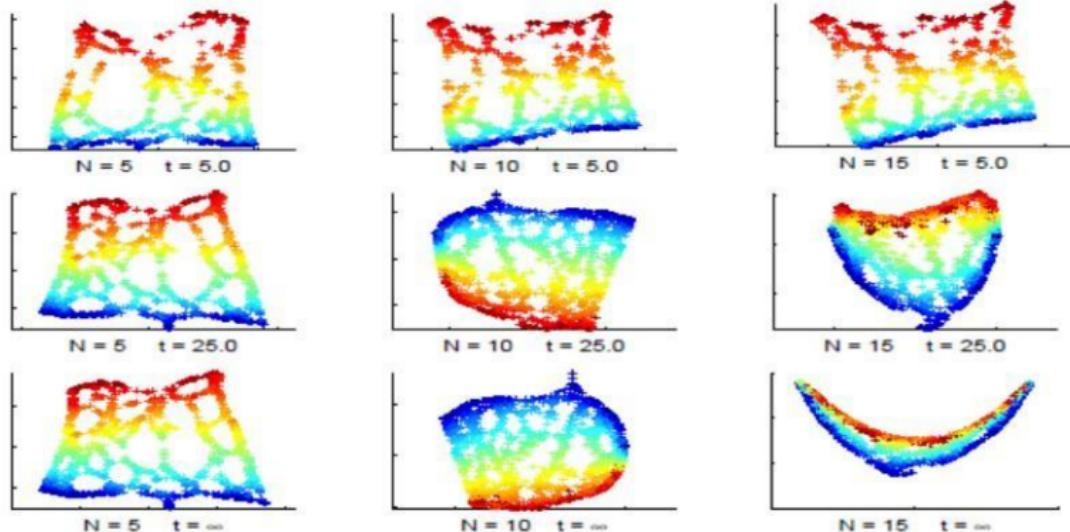
- Swiss roll



2000 random data points on the manifold

# Laplacian Eigenmap Example

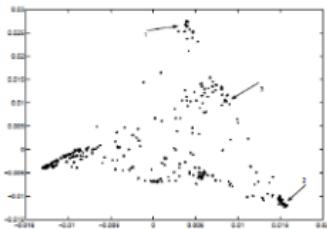
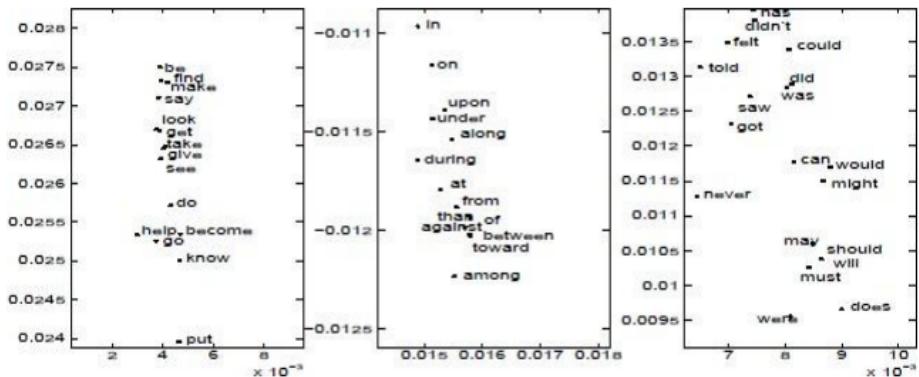
- 2D embedding of the swiss roll



Free parameters,  $N$  and  $t$ .  $N$  = Number of neighbors,  $t$  = Heat kernel parameter

# Laplacian Eigenmap Example

- 300 most frequent words from Brown corpus
- Each word is represented by a 600 dimensional vector
- Laplacian Eigenmap with  $N = 14$ ,  $t = \inf$



Framgments labeled by arrows, from left to right. The first is exclusively infinitives of verbs, the second contains prepositions and the third mostly modal and auxiliariy verbs

# Summary

- Isomap, LLE and Laplacian Eigenmap: non-linear dimensionality reduction technique
- Useful for learning manifolds, understanding low dimensional data embedded in high dimensional space.
- Linear dimensionality reduction technique (PCA, SVD) fails for this type of data.
- All three can preserve local geometry (inter-point relationships)

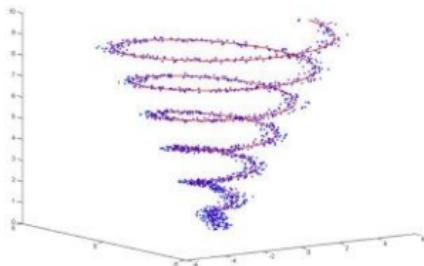
# References and acknowledgments

- Roweis, S. T. & Saul, L. K. (2000), ‘Locally linear embedding’, Science 290, 2323–2326.
- Tenenbaum, J. B., de Silva, V. & Langford, J. C. (2000), ‘A global geometric framework for nonlinear dimensionality reduction’, Science 290, 2319–2323.
- <http://www.cs.nyu.edu/~roweis/lle/>
- <http://isomap.stanford.edu/>
- [http://en.wikipedia.org/wiki/Nonlinear\\_dimensionality\\_reduction](http://en.wikipedia.org/wiki/Nonlinear_dimensionality_reduction)
- [http://web.mit.edu/6.454/www/www\\_fall\\_2003/ihler/slides.pdf](http://web.mit.edu/6.454/www/www_fall_2003/ihler/slides.pdf)
- <http://cseweb.ucsd.edu/~saul/teaching/cse291s07/laplacian.pdf>
- [www.public.asu.edu/~jye02/CLASSES/Spring.../Lec19-Isomap.ppt](http://www.public.asu.edu/~jye02/CLASSES/Spring.../Lec19-Isomap.ppt)
- [www.public.asu.edu/~jye02/CLASSES/Spring.../Lec20-LLE.ppt](http://www.public.asu.edu/~jye02/CLASSES/Spring.../Lec20-LLE.ppt)
- [www.public.asu.edu/~jye02/CLASSES/.../Lec21-LaplacianEig.ppt](http://www.public.asu.edu/~jye02/CLASSES/.../Lec21-LaplacianEig.ppt)

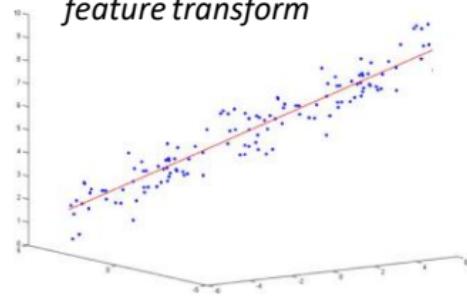
# Kernel-PCA (k-PCA)

Can we kernelize PCA to implicitly do a linear dimension reduction (ie PCA) in a non-linear feature space?

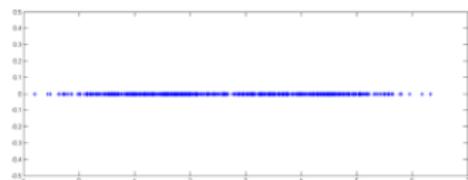
*High dimensional  
manifold data*



*Explicit non-linear  
feature transform*



**Kernel-PCA  
(done implicitly)**



PCA

# Kernel-PCA (k-PCA)

How to kernelize PCA?

Observation:

PCA directions are the principal eigenvectors of the (centered) matrix  $XX^T$   
But... we want the inner product  $X^T X$

Let  $X = USV^T$

- Since  $XX^T = US^2U^T$ , we are interested in finding how any datapoint  $x$  is projected onto the vectors  $U$ , ie the projection  $U^T x$

If data is in a feature space  $\phi(X)$

- Problem: cannot compute  $U$  (no efficient way to compute outer product)
- If dot product can be computed efficiently, then easy to compute  $V$  and  $S$
- Since  $U = \phi(X) V S^{-1}$
- For any datapoint  $\phi(x)$ , its projection  $U^T \phi(x) = S^{-1} V \phi(X) . \phi(x)$

*can be computed efficiently!*

# Kernel-PCA (k-PCA)

We can view many of the manifold embedding methods as a special case of kernel PCA (with specific choice of the kernel)!

- PCA:  $K = X^T X$  (linear kernel)
- Classical-MDS:  $K = -\frac{1}{2} HD^{eucl}H$ , where  $H$  is the centering matrix  
*(Euclidean) distance gets converted into inner product matrix K*
- Isomap:  $K = -\frac{1}{2} HD^{geodesic}H$   
*Basically MDS with geodesic distances*
- Locally Linear Embedding:  $A := (I - W)(I - W)^T$ ;  $K = L^{-1}$  or  $K = (\lambda_{\max} I - A)$   
*W are the learned locally linear weights from LLE*
- Laplacian Eigenmap:  $K = L^{-1}$  or  $K = (\lambda_{\max} I - L)$   
*L is the graph Laplacian used in LE*

# Kernel-PCA (k-PCA)

## Observation

- Many manifold embedding methods assume a **specific kernel form**

Idea: why don't we directly ***learn the kernel*** which yields a good embedding?

Maximum Variance Unfolding (MVU) (aka Semi-Definite Embedding (SDE))

# Maximum Variance Unfolding (MVU)

In order to find an optimal kernel for non-linear embedding, we need to define what we want the embedding to do.

Goal: Find a low-dim. map that preserves the “**local geometry**”  
local geometry = distances between neighboring points!

So, how to learn a kernel that preserve the local distances?

For points  $x_i$  and  $x_j$  in any neighborhood

$$\begin{aligned} \|x_i - x_j\|^2 &= \|\phi(x_i) - \phi(x_j)\|^2 \\ &= K_{ii} + K_{jj} - 2K_{ij} \end{aligned}$$

- $K$  needs to be PSD
- $\sum_{ij} K_{ij} = 0$

*Can formulate the optimization as follows...*

# Maximum Variance Unfolding (MVU)

[Weinberger and Saul '04]

$$\text{Maximize}_K \quad \text{tr}(K) \quad \text{Maximize the spread of points}$$

Constraints:

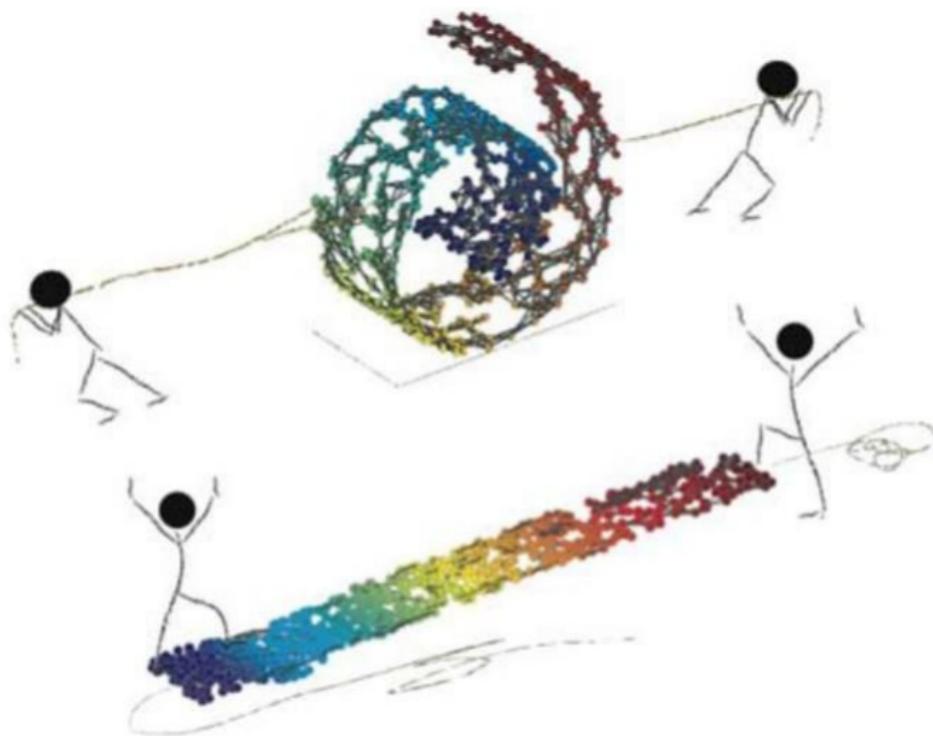
$$K_{ii} + K_{jj} - 2K_{ij} - \|x_i - x_j\|^2 \text{ if } i \text{ and } j \text{ are neighbors.}$$

$$\sum_{ij} K_{ij} = 0$$

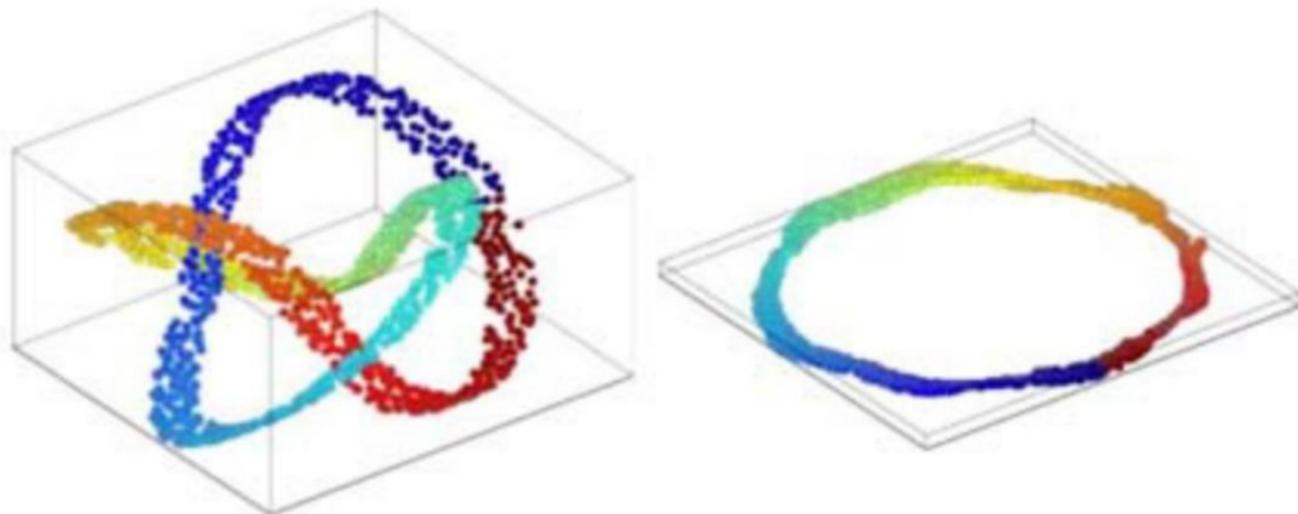
$$K \succeq 0$$

Can be solved by any SDP solver

# Maximum Variance Unfolding (MVU)



# Maximum Variance Unfolding (MVU)



# Stochastic Neighbor Embedding (SNE)

Goal: Find a low-dim. map that preserves the “**local geometry**”

local geometry = **similarity** between points in local neighborhoods

*Similar goal to LE, BUT a  
drastically different solution!*

Idea:

Model the neighborhood structure/information as a probability distribution, then find a low-dimensional mapping that matches the same distribution!

Notation:

- $x_1, \dots, x_n$  given high dim. data (given)
- $y_1, \dots, y_n$  mapped low dim. Representation (to be learned)
- $p_{j|i}$  = probability of  $x_j$  being the neighbor of  $x_i$  (computed from data)
- $q_{j|i}$  = probability of  $y_j$  being the neighbor of  $y_i$  (to be matched to  $p_{j|i}$ )

# Stochastic Neighbor Embedding

[Hinton and Roweis '03]

Stochastic Neighbor Embedding approach:

Probability  
model for high-  
dim input data

$$p_{j|i} = \frac{\exp(-\|x_i - x_j\|^2 / 2\tau_i^2)}{\sum_{k \neq i} \exp(-\|x_i - x_k\|^2 / 2\tau_i^2)}$$

Meta parameter controlling  
the neighborhood size

Probability model  
for low-dim  
mapped data

$$q_{j|i} = \frac{\exp(-\|y_i - y_j\|^2)}{\sum_{k \neq i} \exp(-\|y_i - y_k\|^2)}$$

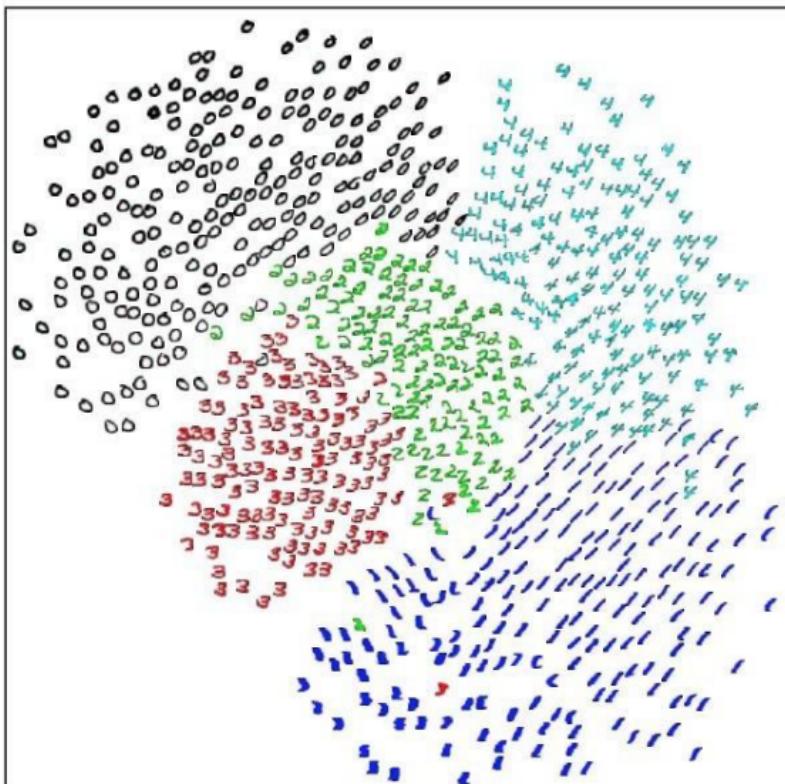
y's are the variables  
that need to be learned

Key optimization: Maximize the similarity between the distributions

$$\text{minimize}_y : \sum_i KL(P_i || Q_i) = \sum_i \sum_j p_{j|i} \log \frac{p_{j|i}}{q_{j|i}}$$

Highly non-convex, just do gradient descent  
and settle with the local optimal solution

# Stochastic Neighbor Embedding



The individual class clusters  
are well all together producing  
an effective visualization

But the clusters are  
NOT well separated

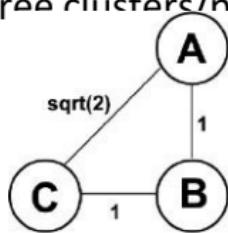
The issue: “crowding  
problem”

# t-distributed Stochastic Neighbor Embedding

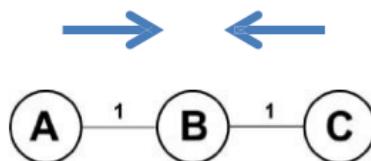
The crowding problem:

High dimensional data is being crammed into a low dimensional space, to match the probabilities, the clusters can “crowd” together

Consider three clusters/points A, B, C



Organization in high dimensions



Organization in low dimensions

Because of the gaussian-type neighborhood structure in low dimensions, large distance between A and C will be **penalized** a lot causing them to be mapped close (ie crowd) to each other

# t-distributed Stochastic Neighbor Embedding

Solution to the crowding problem

[Van der Maaten and Hinton  
'08]

Idea: instead of using a **thin-tailed** Gaussian in the lower dimensions, we can use a **heavier-tailed** distribution, e.g. student's t-distribution!

$$p_{ij} = \frac{p_{j|i} + p_{i|j}}{2n}$$

Symmetrize the high dimensional neighborhood distribution

$$q_{ij} = \frac{(1 + \|y_i - y_j\|^2)^{-1}}{\sum_{k \neq l} (1 + \|y_k - y_l\|^2)^{-1}}$$

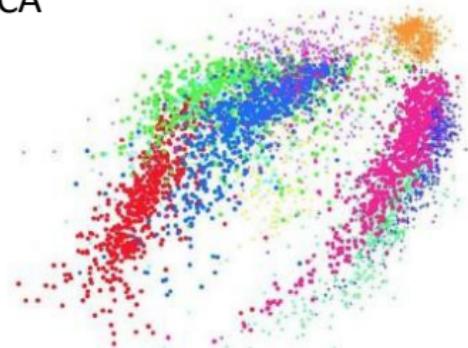
Use the heavier tailed student's t-distribution

Final optimization:

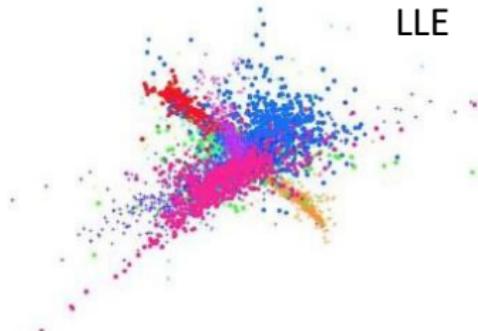
$$\text{minimize}_y \quad \sum_i KL(P_i || Q_i) = \sum_{i=1}^n \sum_{j=1}^n p_{ij} \log \frac{p_{ij}}{q_{ij}}$$

# t-SNE

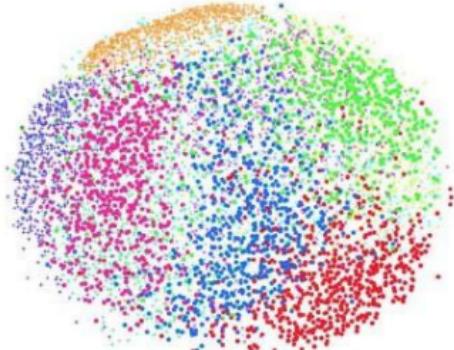
PCA



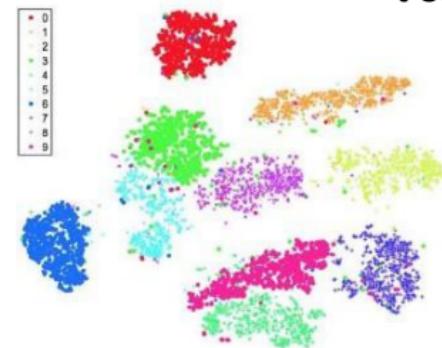
LLE



Sammon mapping

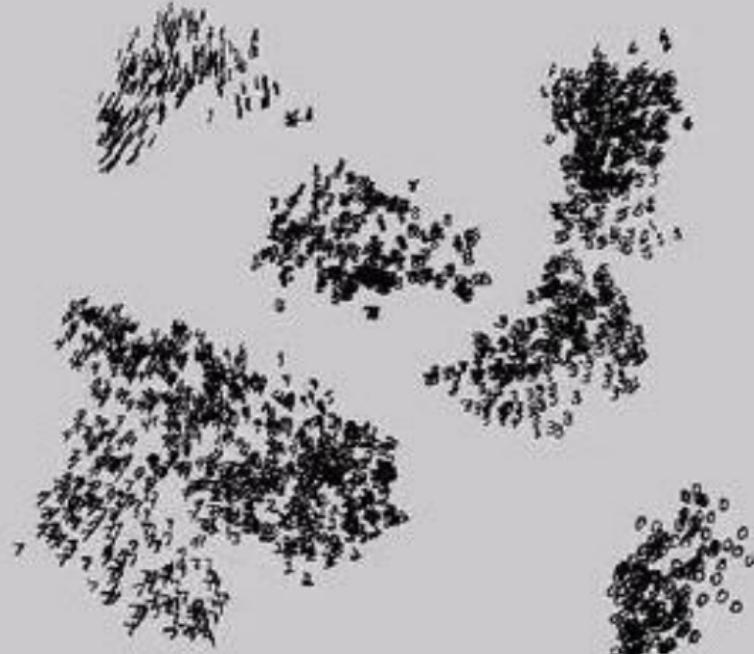


t-SNE



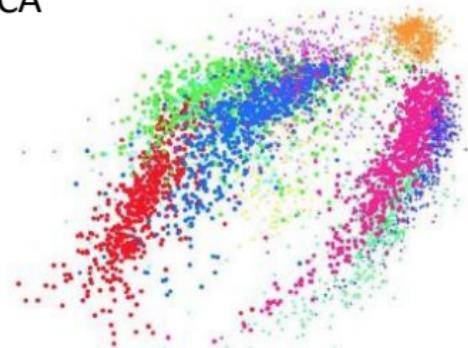
- 0
- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9

# t-SNE

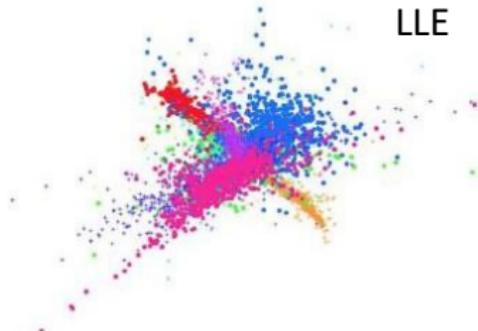


# t-SNE

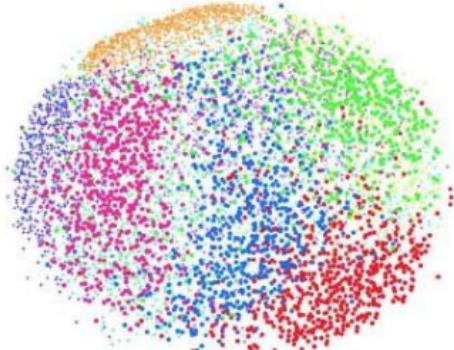
PCA



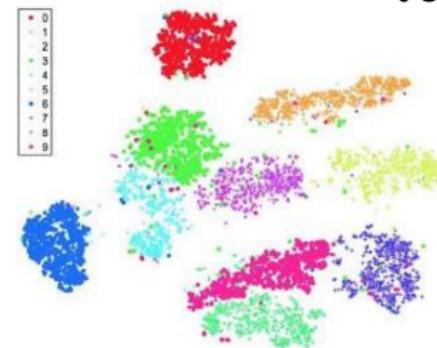
LLE



Sammon mapping



t-SNE



- 0
- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9

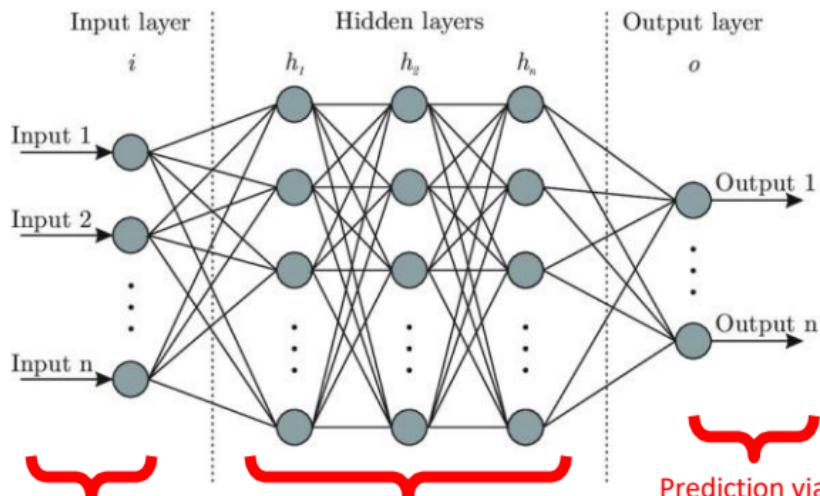
# t-SNE thoughts

- Remarkably effective in visualizing cluster structure in data
- Arguably **the best** ultra low-dimensional technique that “just works”!
  - A few contenders are popping up (eg. UMAP)
    - [McInnes, Healy and Melville '18]
- Tends to cluster even when there may not be any clusters!
  - Can result in false cluster discovery
    - [Im, Verma and Branson '18]
- Recent results have quantified the sufficient conditions needed for t-SNE to be provably successful in revealing the cluster structure in data.
  - [Lindermann and Steinerberger'17, Arora et al. '18]

# Auto-Encoders

Can we use Neural Networks to do (non-linear) dimension reduction?

## Neural Networks



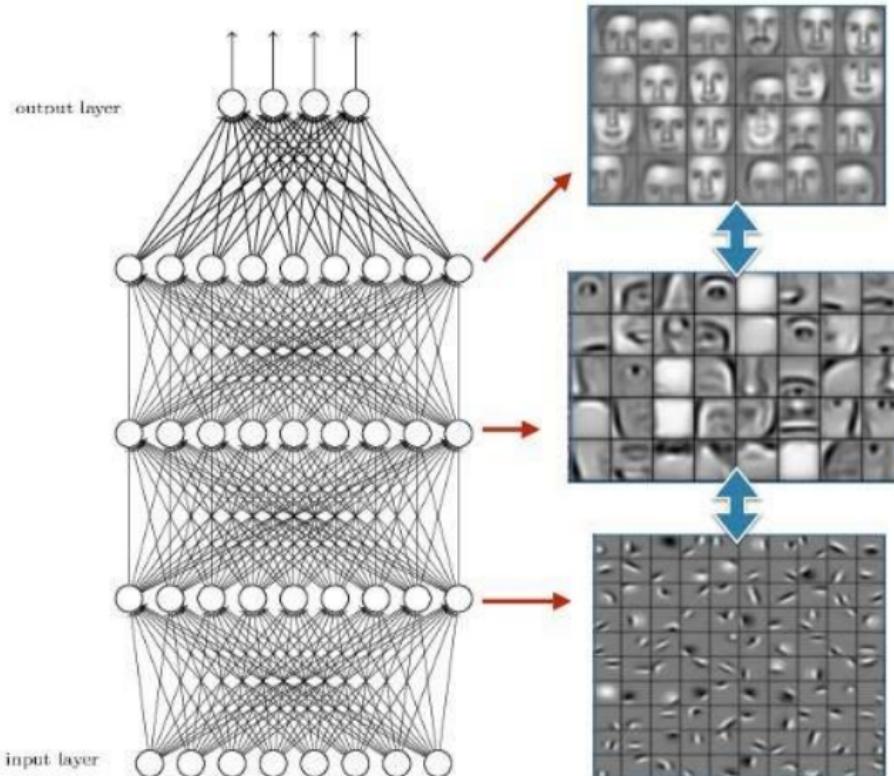
Original representation

Each hidden layer changes the data representation

Weight learning is usually done by back-propagating the error signal (generated by the output label)

Prediction via a simple softmax on the data representation according to the last hidden layer

# Auto-Encoders



Hidden Layers

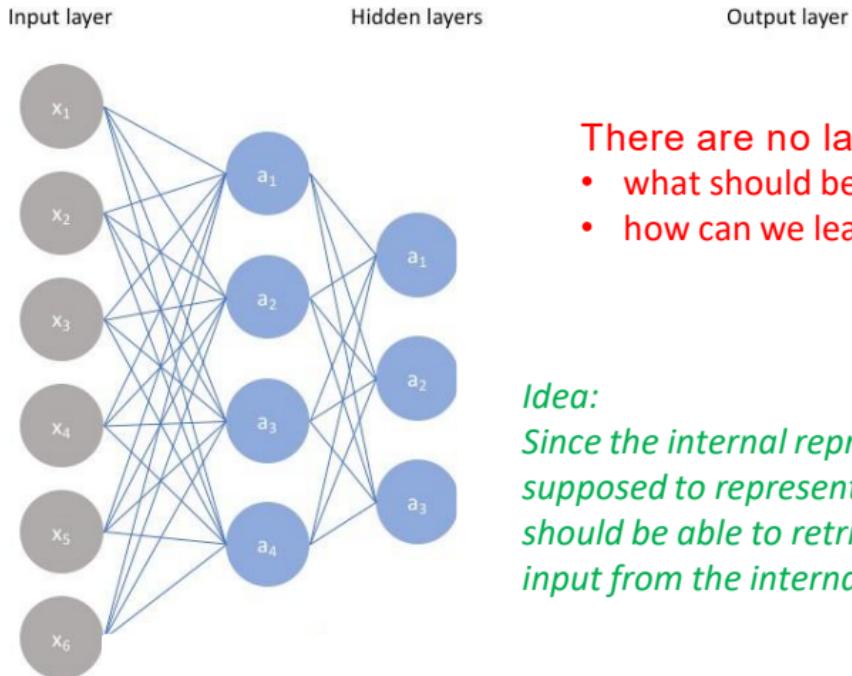
Final/deep hidden layers  
(the representation used for prediction)

Intermediate  
hidden layers

Initial hidden layers

# Auto-Encoders

Idea: To re-represent input data in lower dimensions, we can have the hidden layers of a NN with a narrow width



There are no labels, so...

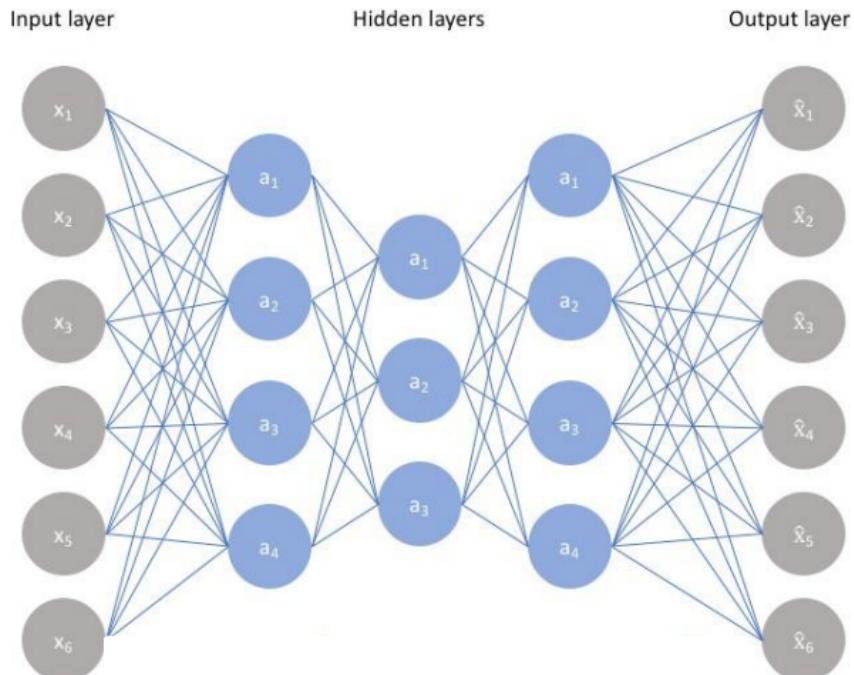
- what should be the output?
- how can we learn the weights?

*Idea:*

*Since the internal representation is supposed to represent the input, we should be able to retrieve/produce the input from the internal representation!*

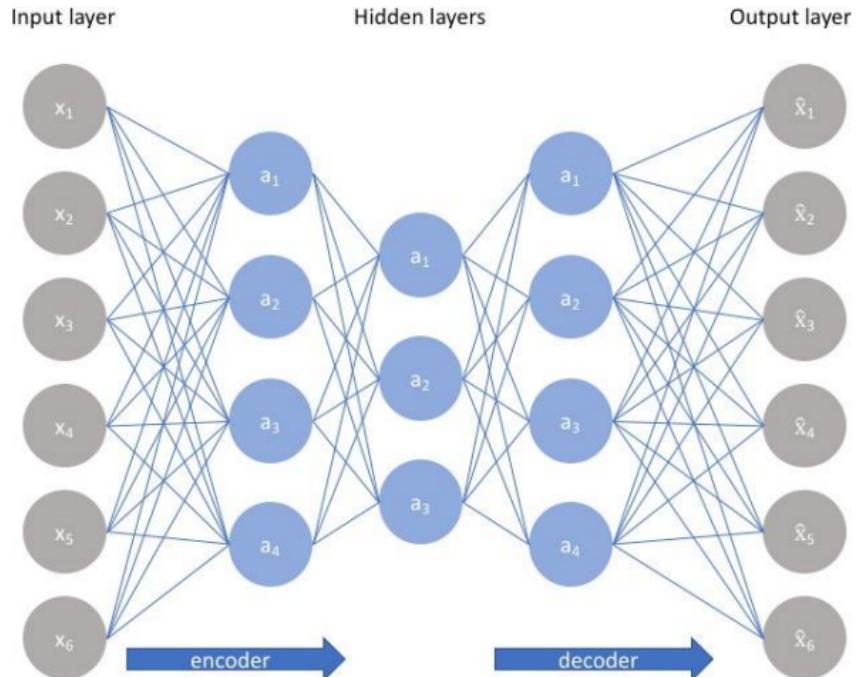
# Auto-Encoders

Idea: To re-represent input data in lower dimensions, we can have the hidden layers of a NN with a narrow width



# Auto-Encoders

*The parameters are learned by backpropagating the error in input reconstruction*



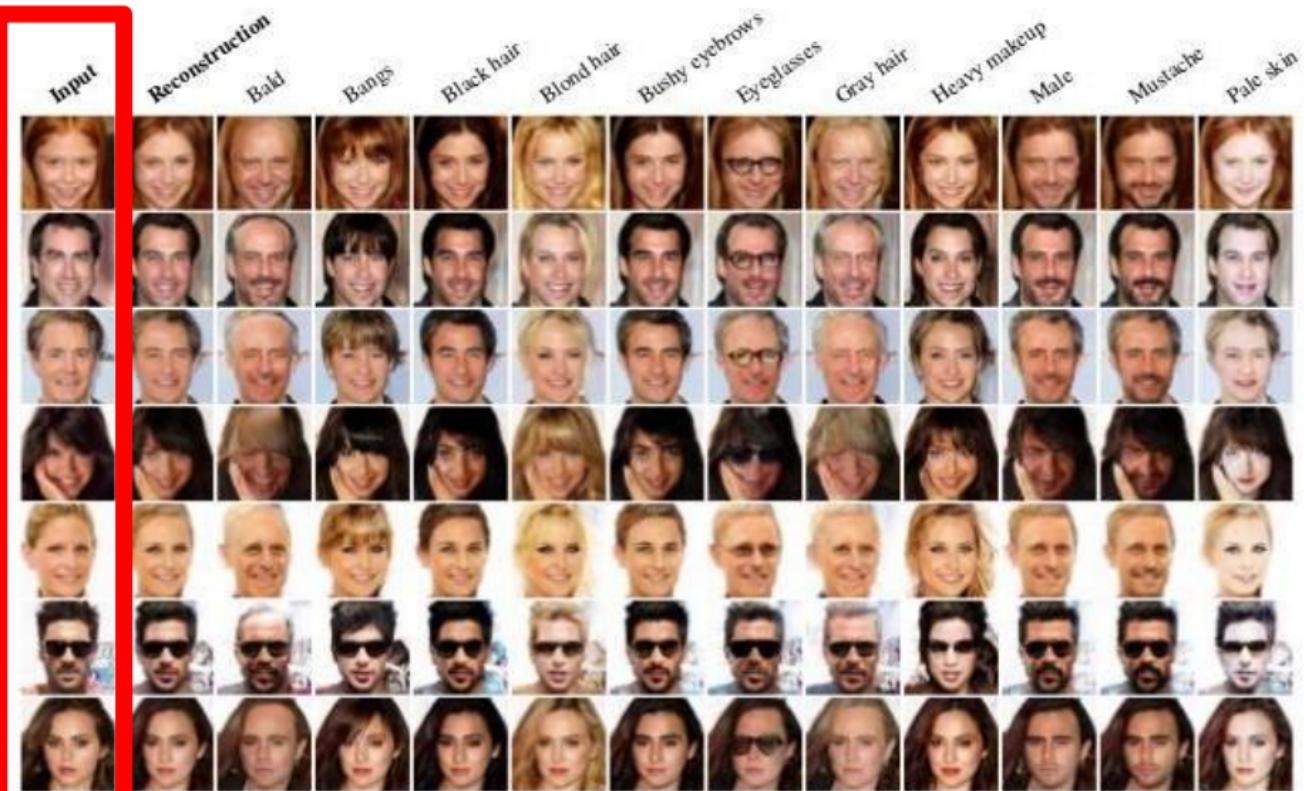
# Auto-Encoders

## Types of Auto-Encoders

- The specific type we saw is called undercomplete auto-encoder  
(hidden layer is of lower dimension than the input layer)
- Overcomplete auto-encoder  
(hidden layer is of higher dimension; it must be regularized)  
can be used for sparse representations
- Denoising auto-encoders (DAE)  
input is deliberately made noisy and the output is required to be a clean representation
- Variational Auto-Encoders (VAE)  
A generative model which can help generate new example datapoints

# VAE

## Potential



## Questions

**Before computing the covariance matrix  $S$  in PCA, you must:**

- A.** Center the data by subtracting the mean from each point
- B.** Standardize each feature to unit variance
- C.**  $\ell_2$ -normalize every vector to unit length
- D.** Remove outliers with k-means

## Questions

**Kernel PCA performs its eigendecomposition on which matrix?**

- A. The raw data matrix  $X$ (size  $N \times D$ )
- B. The kernel matrix  $K$ (size  $N \times N$ )
- C. The graph Laplacian of the data graph
- D. The pairwise distance matrix  $D$

## Questions

**When is classical MDS  
(multidimensional scaling) equivalent  
to PCA?**

- A. Always, regardless of distance
- B. When the dissimilarities are Euclidean distances
- C. Only when the number of points equals the dimensionality
- D. When the kernel is linear

## Questions

**Which method explicitly preserves local neighborhoods by reconstructing each point from its neighbors in both the original and low-dimensional spaces?**

- A. Locally Linear Embedding (LLE)
- B. Isomap
- C. t-SNE
- D. PCA

## Questions

**Compared with SNE, t-SNE differs primarily because it...**

- A. Uses cosine similarity and minimizes JS divergence
- B. Uses a Gaussian in the low-dim space and asymmetric KL
- C. Uses a Student-t distribution in the low-dim space and minimizes a symmetric KL divergence
- D. Uses Earth-Mover distance in a randomized objective

## Questions

K-means++ initialization chooses centers how?

- A. Uniformly at random from  $\mathbb{R}^D$
- B. All from a small random subset of points
- C. First center uniformly at random; each next center is picked from the data with probability proportional to  $D(x)^2$  (squared distance to nearest chosen center)
- D. By running one iteration of standard K-means and taking the largest cluster's points

## Questions

**What is the decision boundary between any two clusters in basic K-means (with Euclidean distance)?**

- A. Circular arc
- B. Quadratic curve
- C. Linear (a hyperplane)
- D. Nonparametric and arbitrary

## Questions

Learning with Prototypes (LwP) using Euclidean distance is equivalent to which classifier?

- A. k-NN with  $k = 1$
- B. Logistic regression with  $L_2$  regularization
- C. A linear classifier with  $w = 2(\mu_+ - \mu_-)$  and  $b = \|\mu_-\|^2 - \|\mu_+\|^2$
- D. Quadratic discriminant analysis

## Questions

**Which matrix appears (inverted) inside the Mahalanobis distance?**

- A. The identity matrix
- B. The kernel matrix
- C. The covariance matrix  $S$
- D. The graph Laplacian

## Questions

**Compared to LwP, which statement about k-NN at test time is most accurate?**

- A. k-NN typically stores all training data and can be slower because it computes distances to all points
- B. k-NN stores only prototypes and is faster
- C. k-NN trains a parametric model with few parameters
- D. k-NN requires no distance computations at test time

## Questions

**For labeled data with overlapping classes, which projection is more suitable and why?**

- A. PCA, because it maximizes total variance
- B. Fisher Linear Discriminant (LDA), because it pushes class means far apart while keeping within-class variance small
- C. Kernel PCA, because it centers the kernel matrix
- D. MDS, because it preserves all pairwise distances