



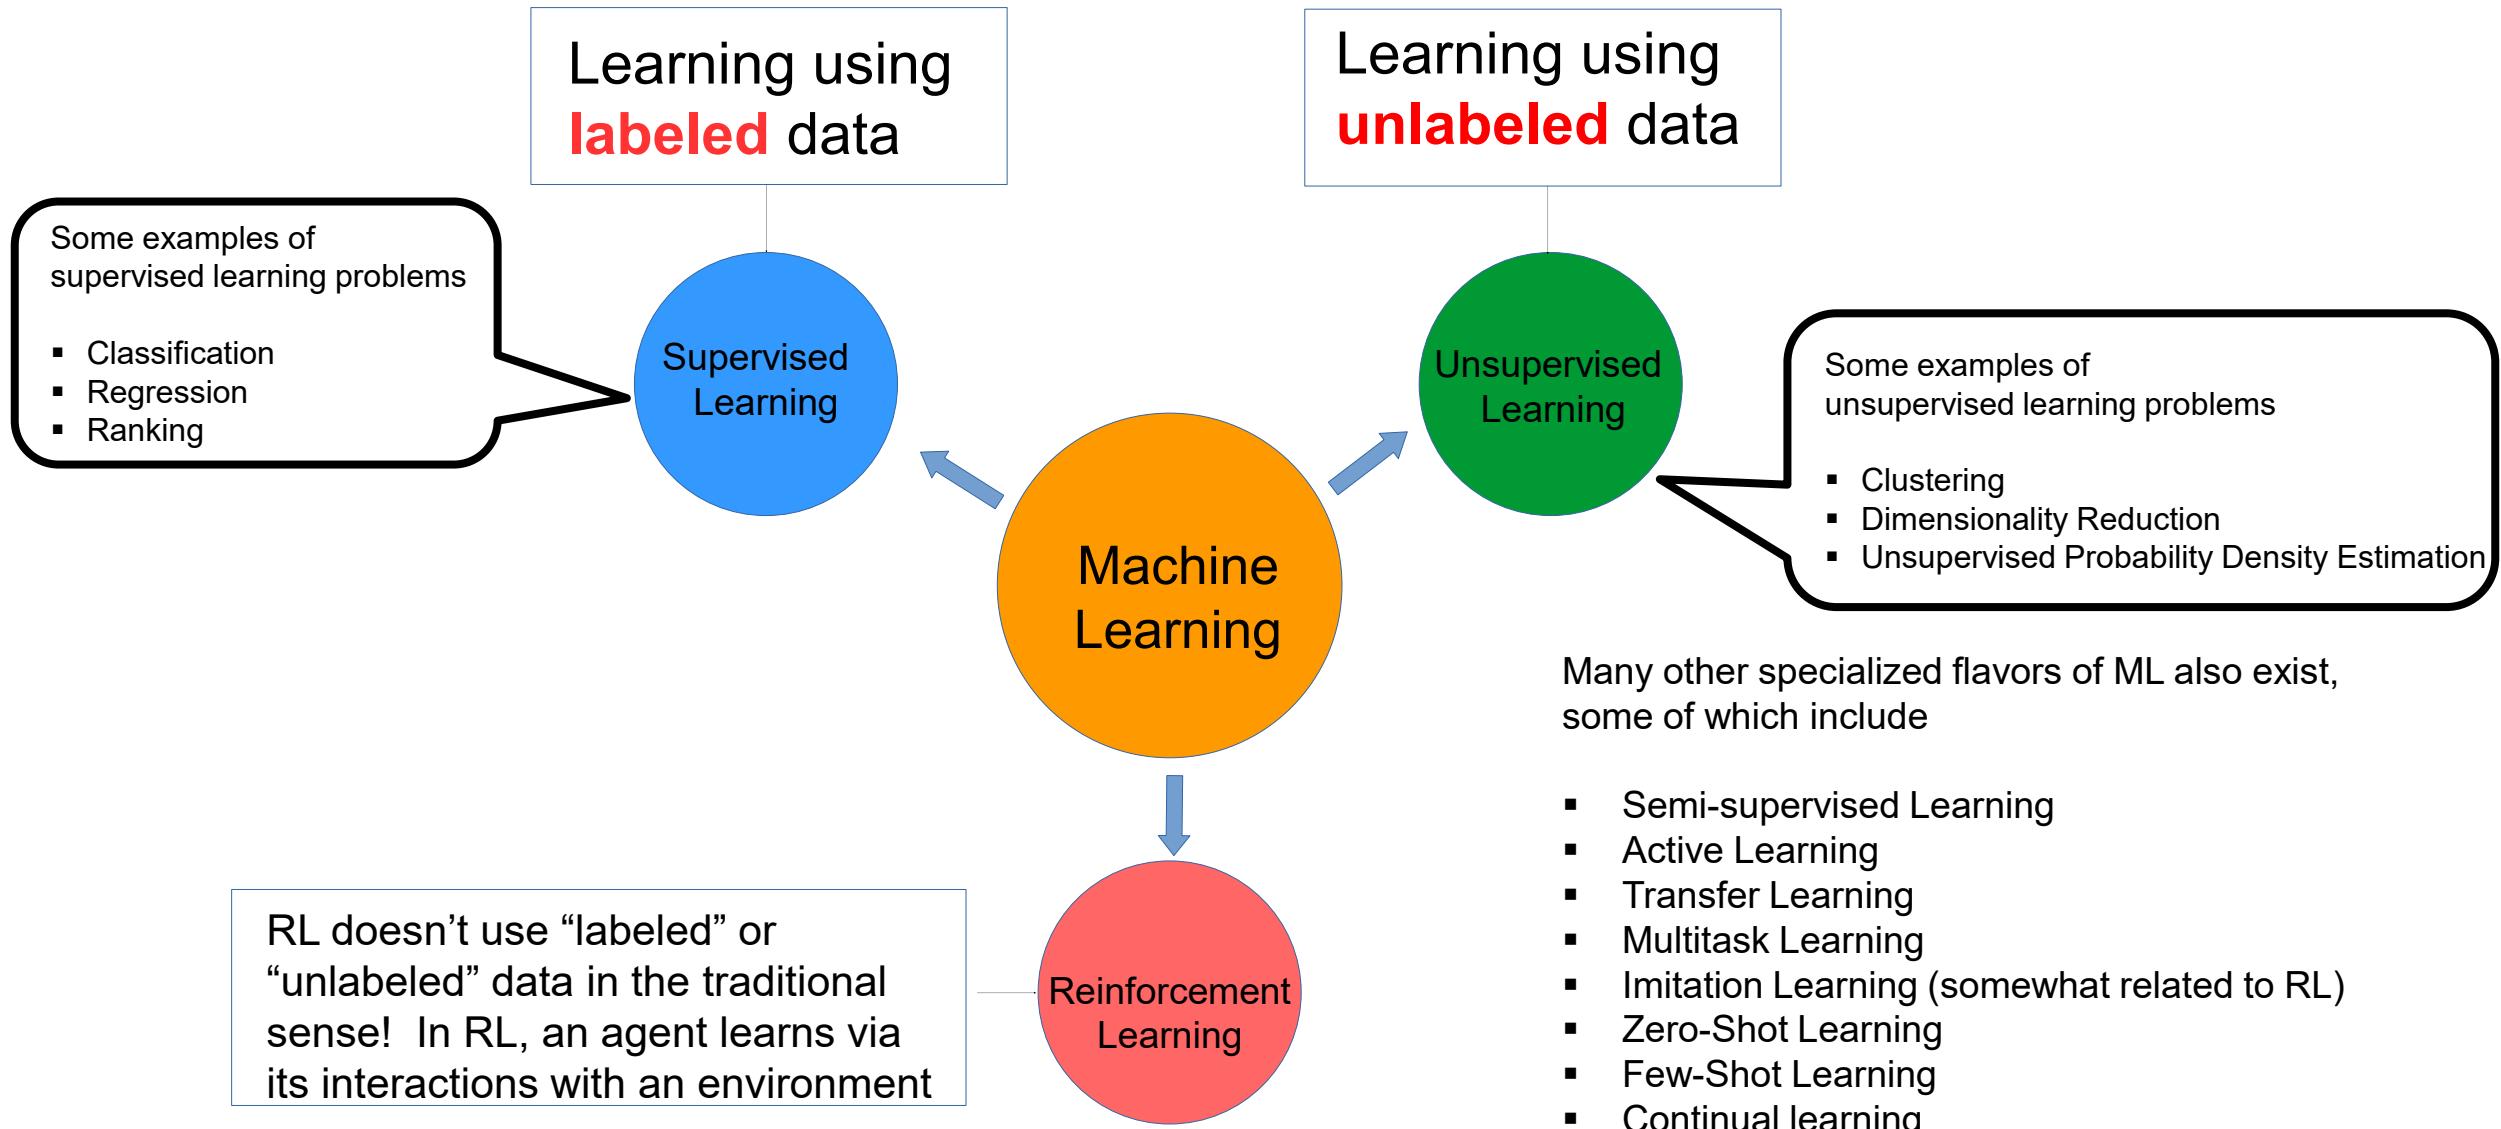
# Working with Data and Features

Dam Quang Tuan

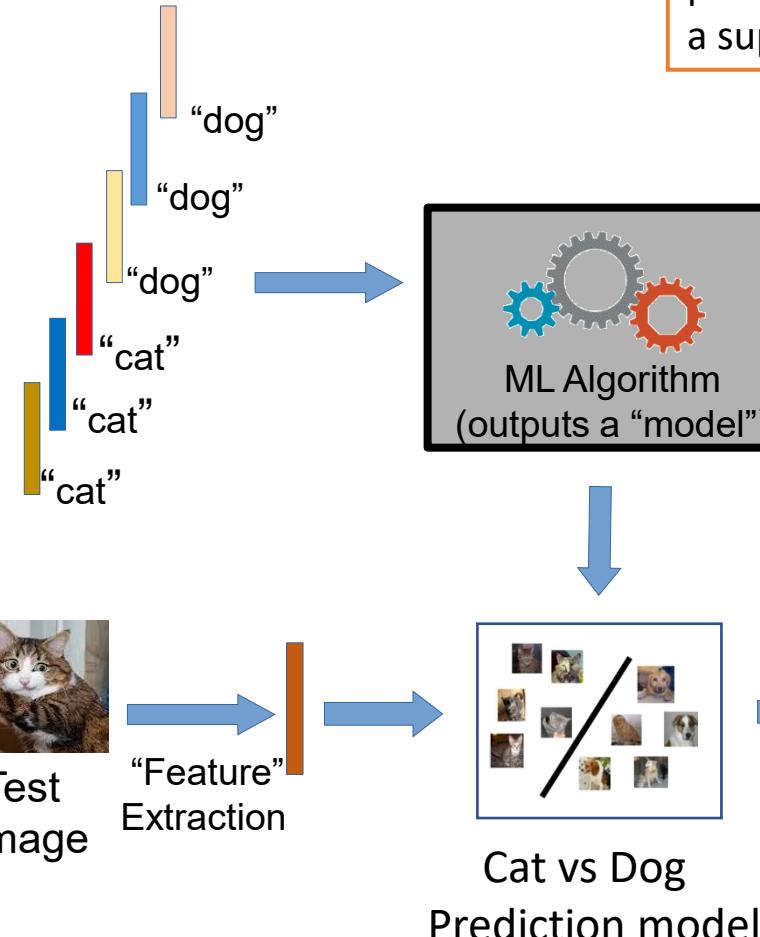
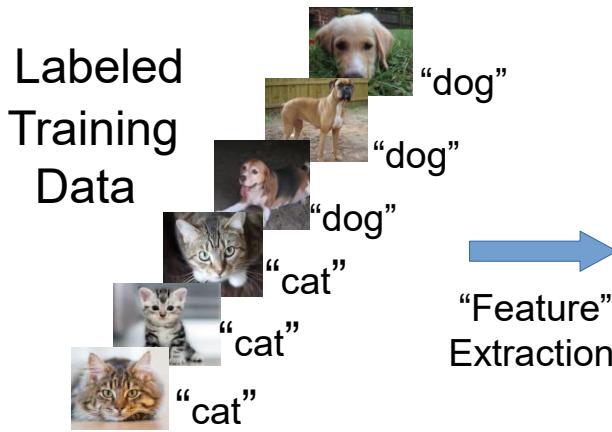
# Plan for today

- Types of ML problems
- Typical workflow of ML problems
- Various perspectives of ML problems
- Data and Features
- Some basic operations of data and features
- People who need a math refresher can use this handy [tutorial](#)

# A Loose Taxonomy of ML



# A Typical Supervised Learning Workflow



Note: This example is for the problem of **binary classification**, a supervised learning problem

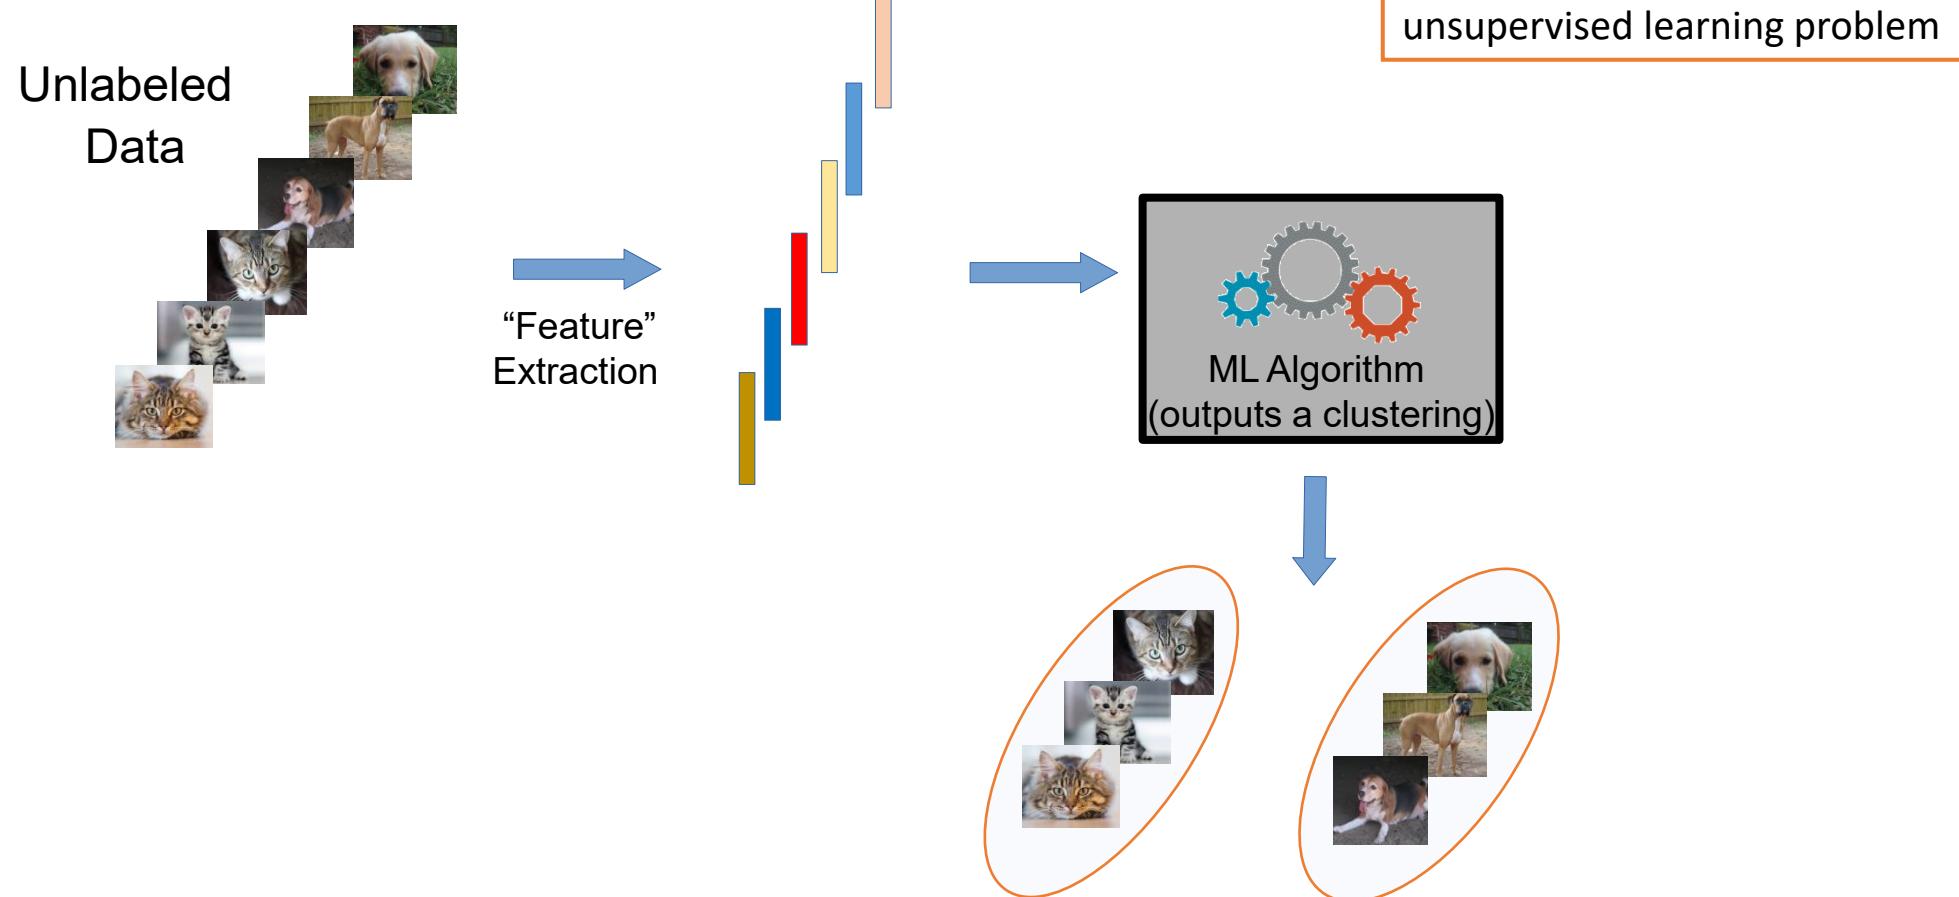


Can you think of a problem you would try to solve using supervised learning?

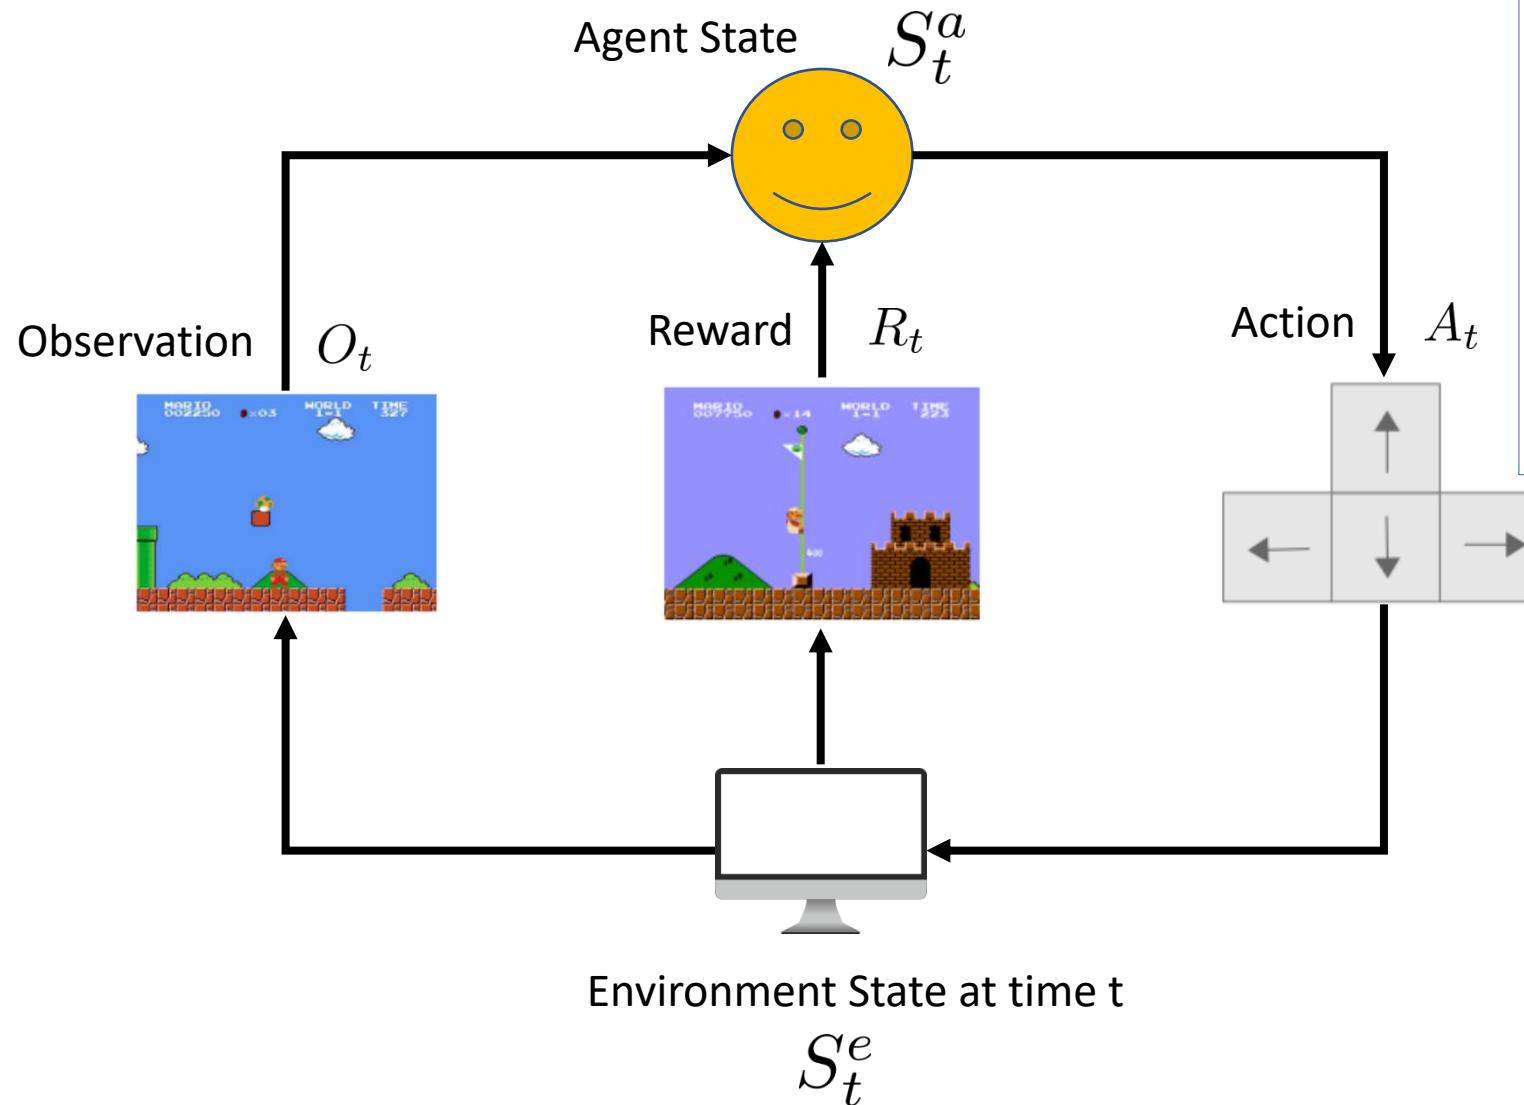
**Feature extraction** converts raw inputs to a **numeric representation** that the ML algo can understand and work with. More on feature extraction later.



# A Typical Unsupervised Learning Workflow



# A Typical Reinforcement Learning Workflow



Wish to teach an agent optimal policy for some task

Agent does the following repeatedly

- Senses/observes the environment
- Takes an action based on its current policy
- Receives a reward for that action
- Updates its policy

Agent's goal is to maximize its overall reward

There IS supervision, not explicit  
(as in Supervised Learning) but  
rather implicit (feedback based)



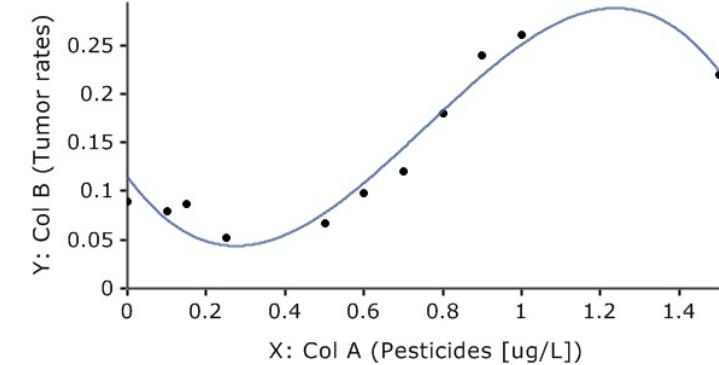
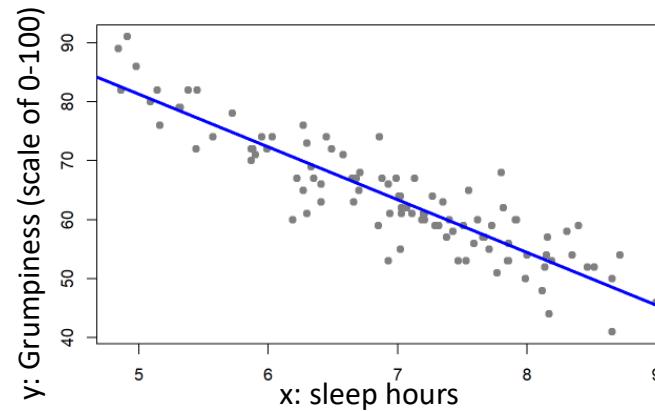
# ML: Some Perspectives

# Geometric Perspective

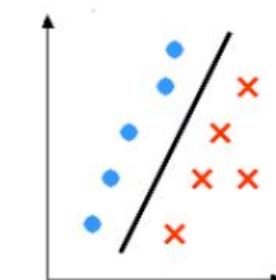
Recall that feature extraction converts inputs into a **numeric representation**

- Basic fact: Inputs in ML problems can often be represented as **points or vectors** in some vector space
- Doing ML on such data can thus be seen from a geometric view

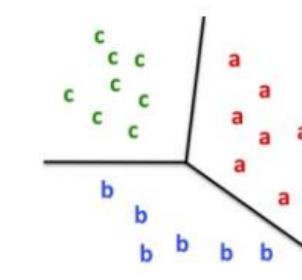
**Regression:** A supervised learning problem. Goal is to model the relationship between input ( $x$ ) and real-valued output ( $y$ ). This is akin to a **line or curve fitting** problem



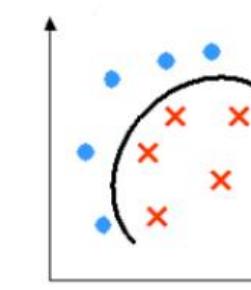
**Classification:** A supervised learning problem. Goal is to learn a to predict which of the two or more classes an input belongs to. Akin to learning **linear/nonlinear separator** for the inputs



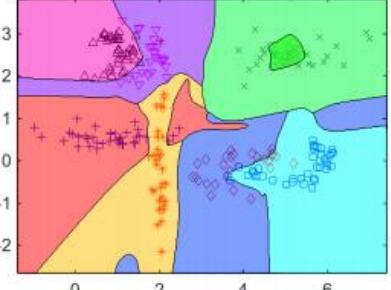
Two-Class (binary)  
Linear Classification



Multi-Class  
Linear Classification



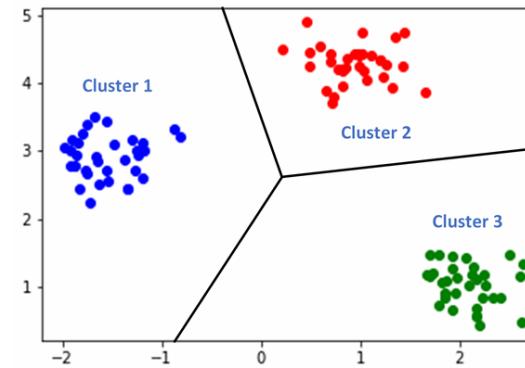
Two-Class (binary)  
Nonlinear Classification



Multi-Class  
Nonlinear Classification

# Geometric Perspective

**Clustering:** An unsupervised learning problem. Goal is to group inputs in a few clusters **based on their similarities with each other**



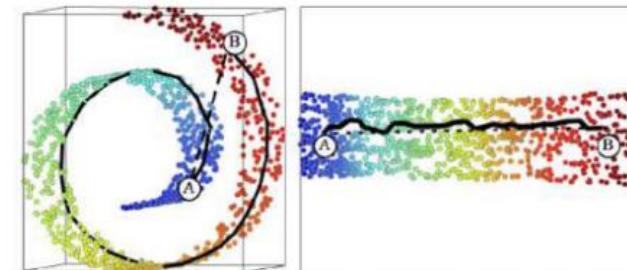
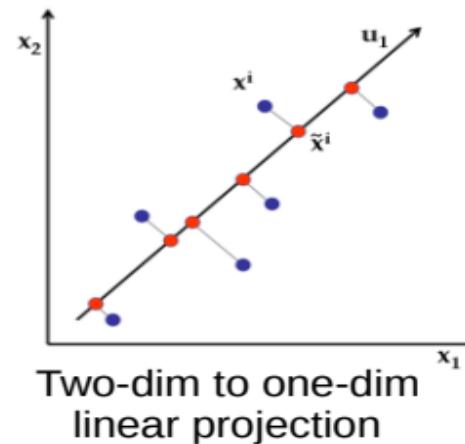
Clustering looks like classification to me. Is there any difference?



Yes. In clustering, we don't know the labels. Goal is to separate them without any labeled "supervision"



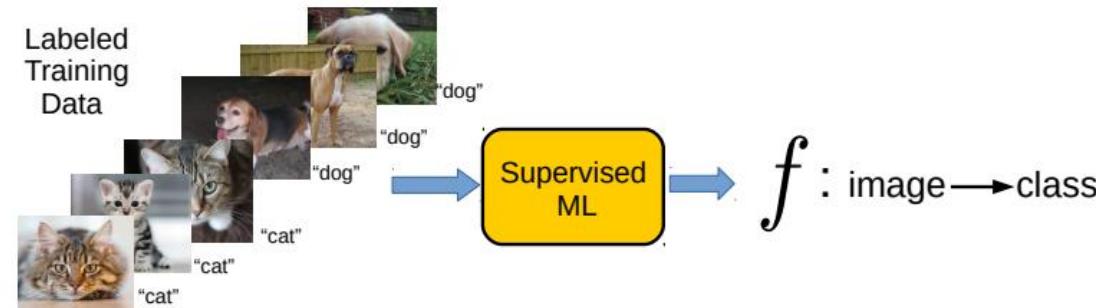
**Dimensionality Reduction:** An unsupervised learning problem. Goal is to **compress the size** of each input without losing much information present in the data



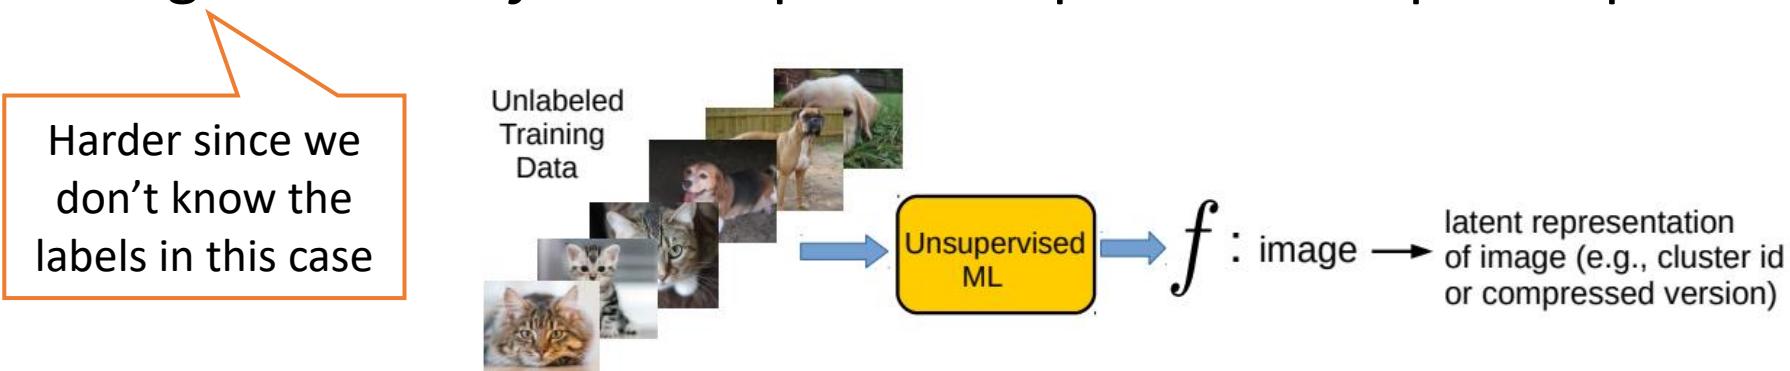
Three-dim to two-dim nonlinear projection (a.k.a. manifold learning)

# Perspective as function approximation

- Supervised Learning (“predict output given input”) can be usually thought of as learning a **function  $f$**  that maps each input to the corresponding output



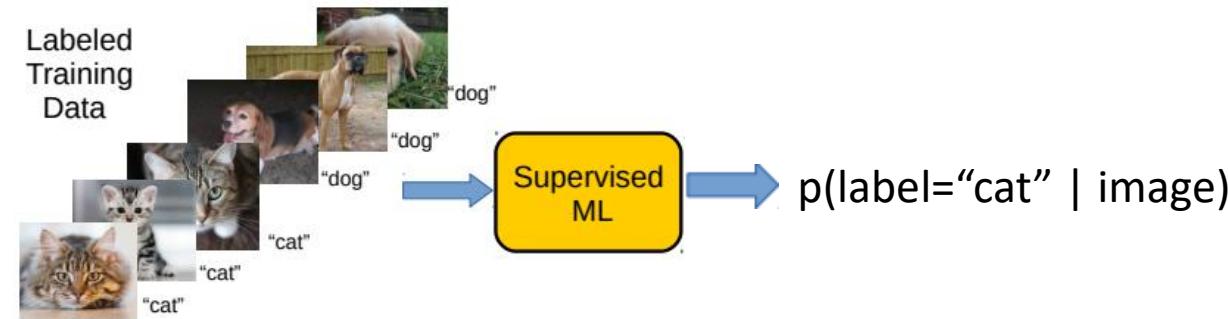
- Unsupervised Learning (“model/compress inputs”) can also be usually thought of as learning a **function  $f$**  that maps each input to a compact representation



- Reinforcement Learning can also be seen as doing function approximation

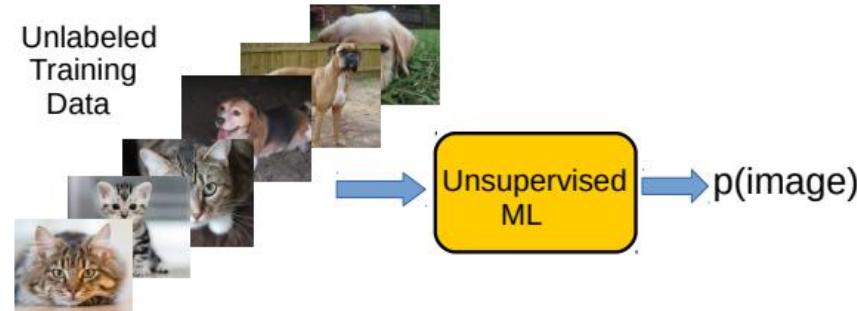
# Perspective as probability estimation

- Supervised Learning (“predict output given input”) can be thought of as estimating the **conditional probability** of each possible output given an input



- Unsupervised Learning (“model/compress inputs”) can be thought of as estimating the **probability density** of the inputs

Harder since we don't know the labels in this case



- Reinforcement Learning can also be seen as estimating probability densities

# Data and Features

# Data and Features

- ML algos require a numeric **feature representation** of the inputs
- Features can be obtained using one of the two approaches
  - Approach 1: Extracting/constructing features manually from raw inputs
  - Approach 2: Learning the features from raw inputs
- Approach 1 is what we will focus on primarily for now
- Approach 2 is what is followed in **Deep Learning** algorithms (will see later)
- Approach 1 is not as powerful as Approach 2 but still used widely

# Example: Feature Extraction for Text Data

- Consider some text data consisting of the following sentences:
  - John likes to watch movies
  - Mary likes movies too
  - John also likes football
- Want to construct a **feature representation** for these sentences
- Here is a “**bag-of-words**” (BoW) feature representation of these sentences

BoW is just one of the many ways of doing feature extraction for text data. Not the most optimal one, and has various flaws (can you think of some?), but often works reasonably well

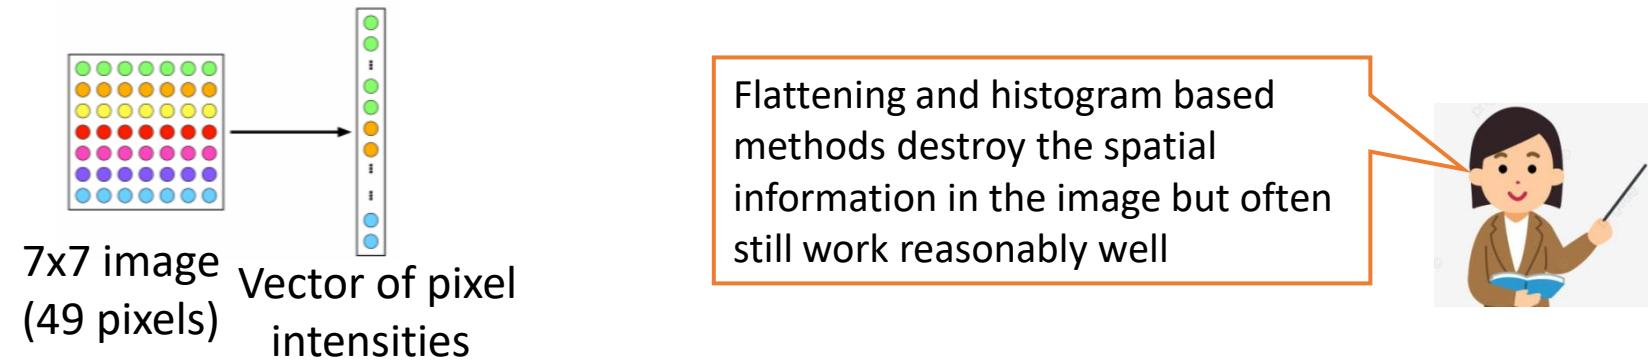


$$\begin{array}{l}
 \text{Sentence 1} \\
 \text{Sentence 2} \\
 \text{Sentence 3}
 \end{array}
 \left( \begin{array}{ccccccccc}
 \text{John} & \text{likes} & \text{to} & \text{watch} & \text{movies} & \text{Mary} & \text{too} & \text{also} & \text{football} \\
 1 & 1 & 1 & 1 & 1 & 0 & 0 & 0 & 0 \\
 0 & 1 & 0 & 0 & 1 & 1 & 1 & 0 & 0 \\
 1 & 1 & 0 & 0 & 0 & 0 & 0 & 1 & 1
 \end{array} \right)$$

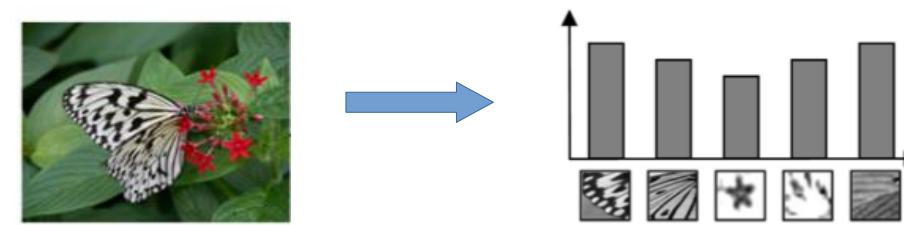
- Each sentence is now represented as a **binary vector** (each feature is a binary value, denoting presence or absence of a word). BoW is also called “**unigram**” rep.

# Example: Feature Extraction for Image Data

- A very simple feature extraction approach for image data is **flattening**



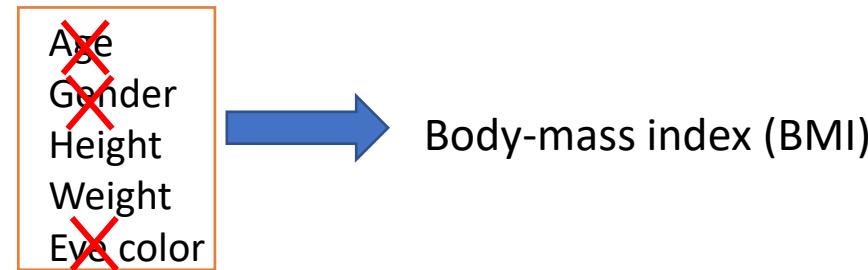
- **Histogram** of visual patterns is another popular feature extr. method for images



- Many other manual feature extraction techniques developed in computer vision and image processing communities (SIFT, HoG, and others)

# Feature Selection

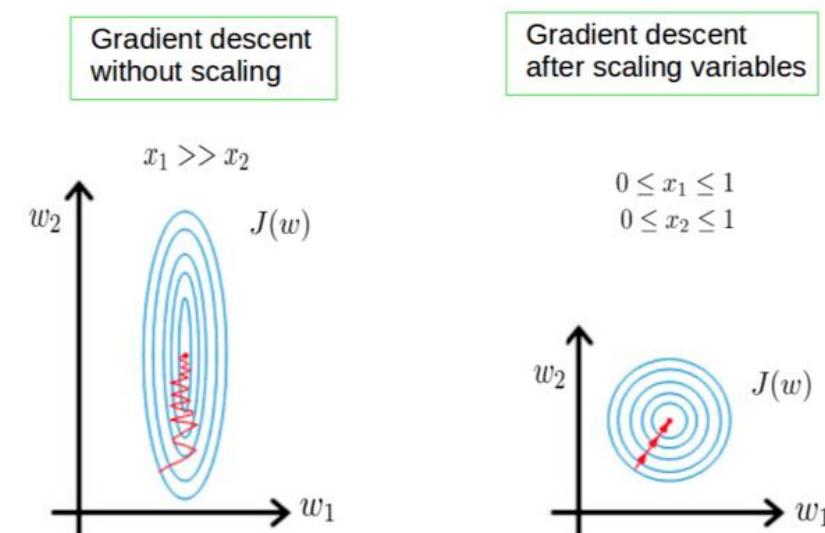
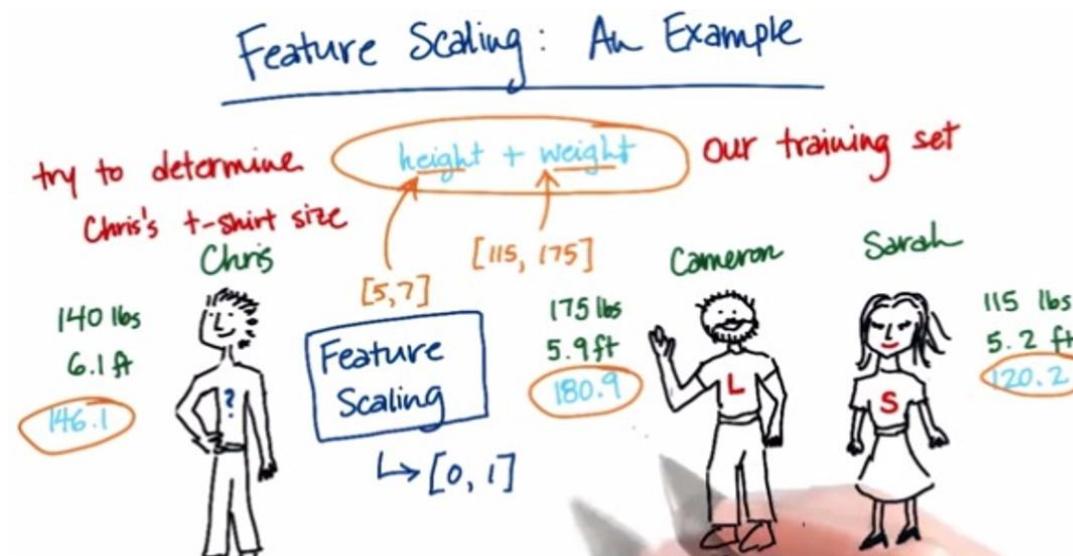
- Not all the extracted features may be relevant for learning the model (some may even confuse the learner)
- **Feature selection** (a step after feature extraction) can be used to identify the features that matter, and discard the others, for more effective learning



- Many techniques exist – some based on intuition, some based on algorithmic principles (will visit feature selection later)
- More common in supervised learning but can also be done for unsup. learning

# Some More Postprocessing: Feature Scaling

- Even after feature selection, the features may not be on the same scale
- This can be problematic when comparing two inputs – features that have larger scales may dominate the result of such comparisons
- Therefore helpful to standardize the features (e.g., by bringing all of them on the same scale such as between 0 to 1)

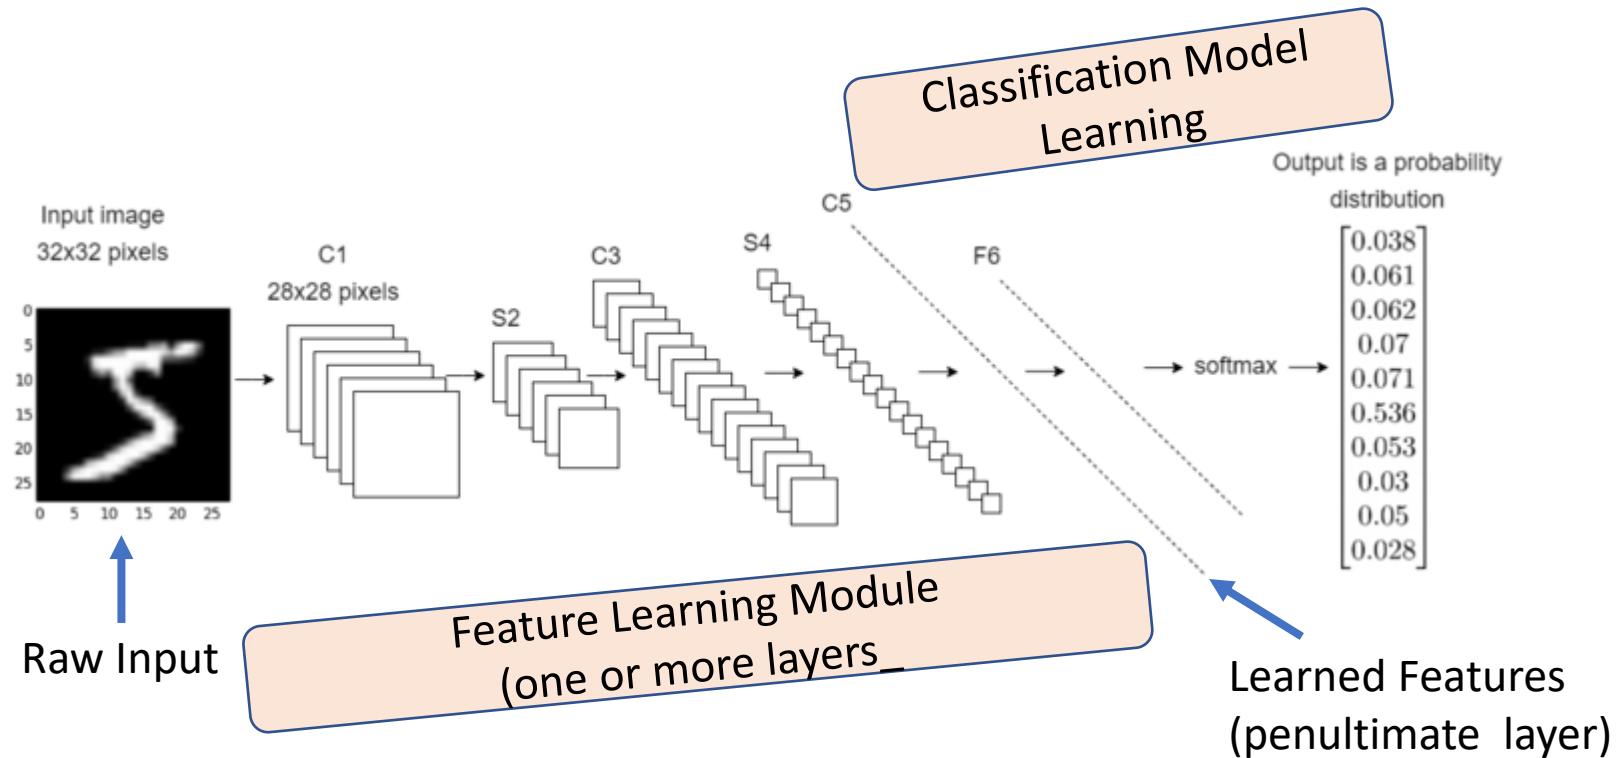


- Also helpful for stabilizing the optimization techniques used in ML algos

# Deep Learning: An End-to-End Approach to ML

Deep Learning = ML with **automated feature learning** from the raw inputs

Feature extraction part is automated via the feature learning module



# Some Notation/Nomenclature/Convention

- Sup. learning requires training data as  $N$  input-output pairs  $\{(\mathbf{x}_n, y_n)\}_{n=1}^N$
- Unsupervised learning requires training data as  $N$  inputs  $\{\mathbf{x}_n\}_{n=1}^N$
- Each input  $\mathbf{x}_n$  is (usually) a vector containing the values of the **features** or **attributes** or **covariates** that encode properties of the it represents, e.g.,
  - For a  $7 \times 7$  image:  $\mathbf{x}_n$  can be a  $49 \times 1$  vector of pixel intensities
- (In sup. Learning) Each  $y_n$  is the **output** or **response** or **label** associated with input  $\mathbf{x}_n$  (and its value is known for the training inputs)
  - Output can be a scalar, a vector of numbers, or even an structured object (more on this later)

RL and other flavors  
of ML problems also  
use similar notation



Size or length of the input  $\mathbf{x}_n$  is  
commonly known as **data/input  
dimensionality** or **feature dimensionality**



# Types of Features and Types of Outputs

- Features as well as outputs can be real-valued, binary, categorical, ordinal, etc.
- **Real-valued:** Pixel intensity, house area, house price, rainfall amount, temperature, etc
- **Binary:** Male/female, adult/non-adult, or any yes/no or present/absent type value
- **Categorical/Discrete:** Zipcode, blood-group, or any “one from a finite many choices” value
- **Ordinal:** Grade (A/B/C etc.) in a course, or any other type where relative values matter
- Often, the features can be of mixed types (some real, some categorical, some ordinal, etc.)

# Some Basic Operations of Inputs

- Assume each input feature vector  $\mathbf{x}_n \in R^D$  to of size  $D$

- Given  $N$  inputs  $\{\mathbf{x}_n\}_{n=1}^N$ , their average or mean can be computed as

$$\boldsymbol{\mu} = \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$$

What does such a  
“mean” represent?



If inputs are all cat images,  
mean vector would represents  
what an “average” cat looks like



- Can compute the Euclidean distance between any pair of inputs  $\mathbf{x}_n$  and  $\mathbf{x}_m$

$$d(\mathbf{x}_n, \mathbf{x}_m) = \|\mathbf{x}_n - \mathbf{x}_m\| = \sqrt{(\mathbf{x}_n - \mathbf{x}_m)^\top (\mathbf{x}_n - \mathbf{x}_m)} = \sqrt{\sum_{d=1}^D (x_{nd} - x_{md})^2}$$

- .. or Euclidean distance between an input  $\mathbf{x}_n$  and the mean  $\boldsymbol{\mu}$  of all inputs

- .. and various other operations that we will look at later..

# Next Class

- Introduction to Supervised Learning
- A simple Supervised Learning algorithm based on computing distances

# References

CS771: Intro to Machine Learning (Fall 2021), Nisheeth Srivastava, IIT Kanpur